

Deployment of a Django Application (Part 1 – Server Setup)

This chapter details the initial setup of a virtual server to host a Django application. We will start by learning about the architecture of deployment and how the frontend and the application servers work together. Then, through a detailed step-by-step walk-through, we will set up an Ubuntu Linux virtual machine that can run on either a PC or macOS. Next, we will study **SSH** and its basics, and we'll use the `ssh` command to connect to our newly set up server. Then, using the `sudo` command over an SSH connection, we will learn how to administer a server remotely. Toward the end of the chapter, we'll shore up our server administration skills by learning how to update and install packages with `apt`. We'll then use `apt` to install **NGINX**, **PostgreSQL**, and **Gunicorn** on our virtual server.

In this chapter, we will cover the following topics:

- Server architecture
- Working with Linux
- Packages and updates

Server architecture

So far, we have been using the Django development server to both execute the Python code for Bookr and serve the static and media files. We mentioned throughout the book that the Django development server is not suitable for use in production. It is not designed to run multiple processes or handle many users, and more importantly, it doesn't run when the `DEBUG` setting is set to `True`.

We will start this chapter by looking at the architecture used in a production web server – that is, where the functions of the Django development server are split and handled by two different applications. There is a *frontend web server* (for example, NGINX, lighttpd, or Apache), which receives the request from the browser, and the *application server* (for example, Gunicorn or uWSGI), which executes the Python code. The frontend web server decides how the request should be handled. If it is a request for a static or media file, then the frontend web server can handle the request itself; it can just read and send the file. If the request is for Python-specific code, then it will be forwarded to the application server to handle. The Django application then parses the URL and other HTTP data, generates an `HttpRequest` object using a view, and sends the response back to the frontend web server. The frontend web server then passes the response back to the browser.

The following figure elaborates on this process:

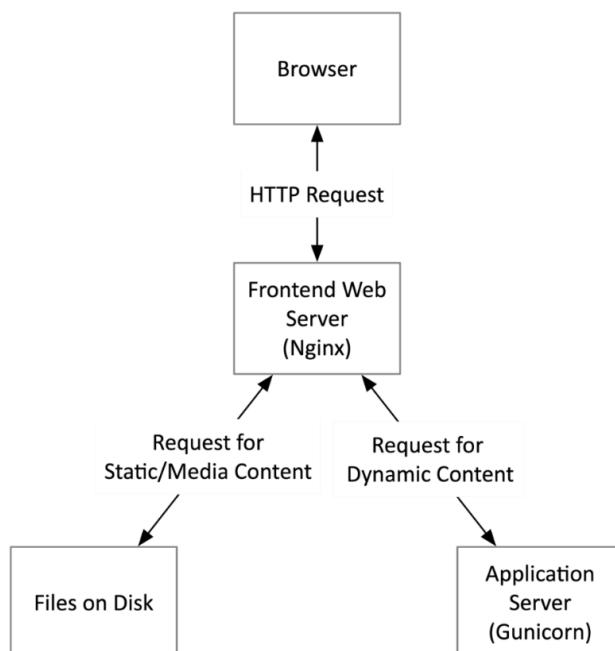


Figure 17.1: Request validation

We will discuss **Gunicorn** and **NGINX** (and explain what they are) in the upcoming sections, titled *NGINX* and *Gunicorn* respectively. While there are plenty of options out there for servers, we're using NGINX and Gunicorn because they're fast, lightweight, and easy to configure. Usually, generating a web page's content using Python (or another programming language) takes much longer than it does to just serve a static or media file off a disk. Therefore, the performance of a website can be increased by adding more backend application servers behind a single frontend server. You can run multiple instances of the application server on the same machine as the frontend server. Due to modern hardware

having multiple CPUs, this can provide speed benefits. For even better performance, you can move the application server onto its own separate server or even run multiple application server instances across different servers. The frontend server can delegate requests to backend servers sequentially or to the one that is currently least busy. The architecture of this would look like the following:

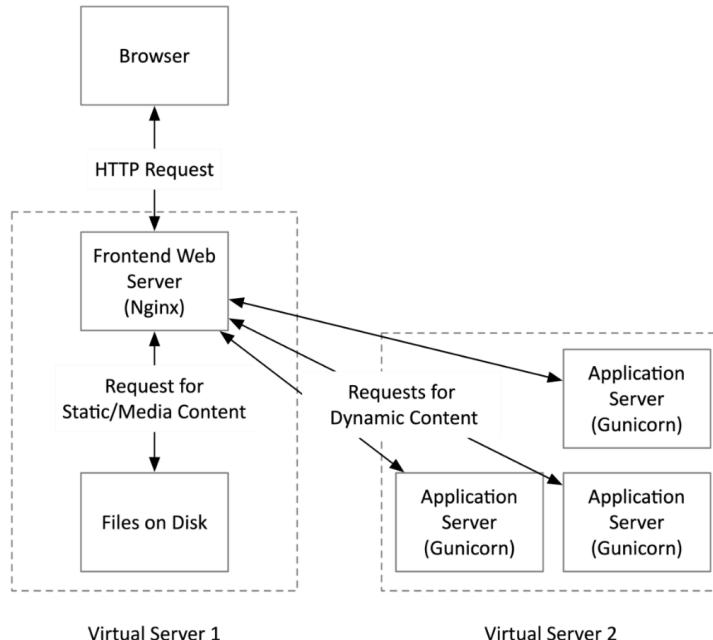


Figure 17.2: Architecture with multiple backends

Figure 17.2 shows how the server architecture is split across multiple servers. A single frontend server, **Virtual Server 1**, stores static files and serves them from its local disk. Requests for dynamic content are forwarded from **Virtual Server 1** to **Virtual Server 2**, which runs multiple **Gunicorn** instances. We could further expand this with multiple virtual servers that follow the same pattern as **Virtual Server 2**.

In this chapter, we will use NGINX as the frontend server and Gunicorn as the backend (or application) server. We'll also switch to using PostgreSQL as the database server instead of SQLite. This is because SQLite is designed just for a single user and is fine for development purposes when you're running Django on your own computer. PostgreSQL is a database server designed for high loads that can scale to thousands of connections.

Gunicorn

Gunicorn is an easy-to-use application server. Its settings can be controlled with command-line flags or a configuration file. In its simplest usage, it can be run like this:

```
gunicorn bookr.wsgi:application
```

This starts Gunicorn using the Django WSGI.

Note

WSGI (often pronounced *wizz-gee*) stands for **Web Server Gateway Interface**. It describes a standard for Python to communicate with web servers. This means any WSGI Python script should be able to communicate with any WSGI server, and any combination of frontend server and application server can be used together, provided they both “speak WSGI.” For example, instead of using Gunicorn, we could use uWSGI. Its configuration would be different, but it serves the same purpose.

We will return to Gunicorn’s configuration in *Chapter 18, Deployment of a Django Application (Part 2 – Configuration and Code Deployment)*.

Note

Prior to this chapter, the code and examples we gave worked on Windows, macOS, or Linux. In this chapter, the installation and setup guides assume an Ubuntu Linux server. This is because the paths to data and config files differ between operating systems. It is possible to run NGINX and Gunicorn on Windows or macOS too, but you will need to research how to install them yourself.

You can use **VMware** or **VirtualBox** to configure a Linux virtual machine running on your computer or use a hosted virtual private server from a provider such as **Amazon Web Services (AWS)**. If you use a virtual machine, you can shut it down when you are not working on it to free up resources on your computer.

I will give instructions on how to set up an Ubuntu Linux virtual machine running inside VirtualBox on your computer, but if you’re more familiar with other virtualization software or have a cloud provider that you would like to use, you can set up a virtual server in that environment instead. Provided you create a virtual server running the same version of Ubuntu, the instructions should be the same after setup is complete.

Throughout this chapter, we will refer to this virtual machine as the “virtual machine.” If a command is to be run on it, we will say “*Run command... on your virtual machine,*” “*On your virtual machine, do...,” and so on.*

A hosted virtual server

A simple way to get your website online is to use a single virtual server that runs all the software your application needs. This will include the frontend web server, the application server, the database server, and your Django code. As your site grows in popularity or requires extra performance or features, you can start splitting out frontend and backend applications to be run on separate servers.

For now, we will discuss the basic steps to get a Django website running on a new virtual machine. Since it will be running on your local machine, it will only be accessible from your computer and to others on your networks (for example, other people in your home or office). The process for setting up Django is the same for a hosted virtual server, so once you are ready to take an application online, you can follow the same steps and deploy it to a machine that the public can access.

Note

There are a number of virtual server providers that you can use to quickly set up a virtual server. These include **AWS** (<https://aws.amazon.com>), **DigitalOcean** (<https://www.digitalocean.com/>), **Linode** (<https://www.linode.com/>), and others. If you already have an account with a virtual server provider like this, you can choose to create a virtual server with them. Once you have the server set up and running, then the instructions should be the same.

Once the virtual server is set up, you will connect to it using SSH. SSH is built into macOS and Linux, but a third-party SSH client should be installed for Windows. We will use PuTTY. We can then start installing the software packages necessary to run our Django website. These will not only be things such as Python but also the supporting software, such as the frontend web server (NGINX) and database server (PostgreSQL).

In *Chapter 18, Deployment of a Django Application (Part 2 – Configuration and Code Deployment)*, after the software is installed, we can create a PostgreSQL database and user. We'll then create the web server directory to hold the static and media files.

Note

During server setup, we will use some basic Linux system administration commands (such as `mkdir`, `chmod`, and `sudo`). If you haven't used these before, don't worry. We will explain them as we use them.

Next, we will need to create a system user on the virtual server that has permission to read, write, and execute the Django project. We will then execute the Gunicorn process as the newly created user.

Note

We use a non-root user for security. If there are security issues in our Django code that expose access to other parts of a system, running a server as a non-root user will limit the access of an attacker only to our application. If running it as root, the attacker could access the entire server. The root user can also start other servers, modify configuration files, and delete data for other users. It is good practice to separate applications to be run by their own users so that if one gets compromised, other applications can't be accessed.

With that in place, we can upload our Django code. This is done over **Secure File Transfer Protocol (SFTP)**. In this book, we will use **FileZilla**, a free, multi-platform SFTP client, but you can use any SFTP client if you are already familiar with another one. Before uploading the files, we should configure the production settings in `settings.py` and make a minor change to `wsgi.py` to enable the use of the **Django configurations** that were added in *Chapter 15, Django Third-Party Libraries*. We should then make sure our library requirements are up to date in `requirements.txt`.

After uploading the Django project, we can create a new virtual environment. We'll then install the requirements listed in `requirements.txt` using `pip3`. We can then run `collectstatic` to copy all the static files into the web server directory. Finally, we'll do our configuration setup by editing/creating configuration files for the server. We'll need to create `systemd` scripts to start Gunicorn and edit the NGINX configuration so that it knows how to load static files and pass on other requests to the Gunicorn instance.

This sounds like a lot of steps, but we'll go through them one by one until Bookr is up and running on our own server.

The first thing to do is to get our own virtual server up and running. In the next exercise, we will download an Ubuntu Linux ISO (a virtual installation disk) and VirtualBox, a free virtualization program that runs on many different platforms. If you already have a different preferred virtual machine host or cloud provider that you wish to use, you can skip *Exercise 17.01 – VirtualBox installation and virtual machine creation*, *Exercise 17.02 – Ubuntu Linux installation*, and *Exercise 17.03 – VirtualBox networking configuration*.

Exercise 17.01 – VirtualBox installation and virtual machine creation

In this exercise, we will create a virtual machine that is ready for Ubuntu setup. First, we will download an Ubuntu Linux installation ISO and VirtualBox installation software. We'll set up VirtualBox and a new virtual machine and then configure it, ready to install and run Ubuntu:

1. In a web browser, go to the Ubuntu Linux server download site at <https://ubuntu.com/download/server>. Under the **Ubuntu Server 22.04.1 LTS** header, click the **Download** button. While version 22.04 is not the most recent version, it is a **Long-Term Support (LTS)** release, which means that it is mature and stable and will be supported for a long time (until April 2032):

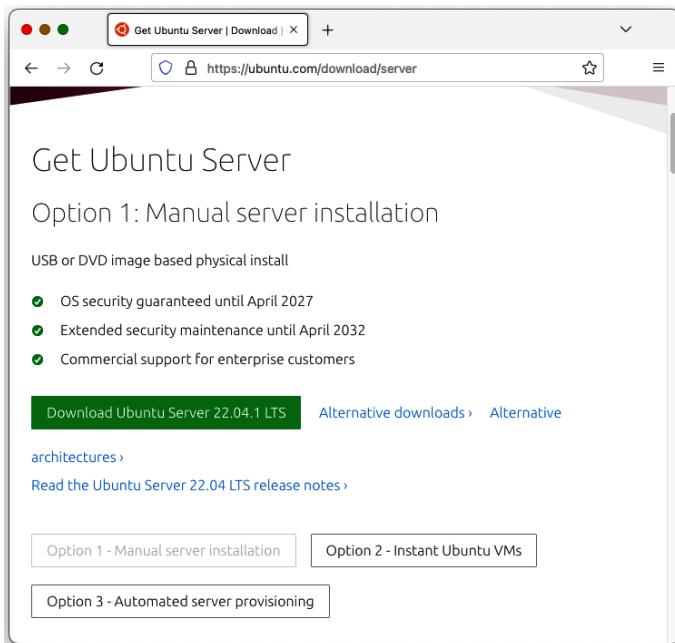


Figure 17.3: The Ubuntu Server 18.04.1 download page

Note

Note that the “patch” portion of the version might change by the time you’re reading this – that is, the version may have moved up to **22.04.2**, **22.04.3**, and so on. If the first two numbers are **22.04**, then this doesn’t matter.

You will be taken to the **Thank You** page, and the ISO will begin downloading. It is around 870 MB, so it might take some time depending on the speed of your internet connection:



[Show all downloads](#)

Figure 17.4: The Ubuntu ISO download prompt in Firefox

2. While the Ubuntu ISO is downloading, you can also download and begin setting up VirtualBox:
 - I. Go to the VirtualBox download page at <https://www.virtualbox.org/wiki/Downloads>.
 - II. Click the download link for your operating system under the **VirtualBox 7.0.4 platform packages** header. Please note again that this version might have changed by the time you read this.
 - III. Click **Windows hosts** if you have a Windows computer, **OS X hosts** if you use macOS, and so on. If you're running Linux as your host operating system, you can click **Linux distributions** to find instructions for your particular Linux distribution:

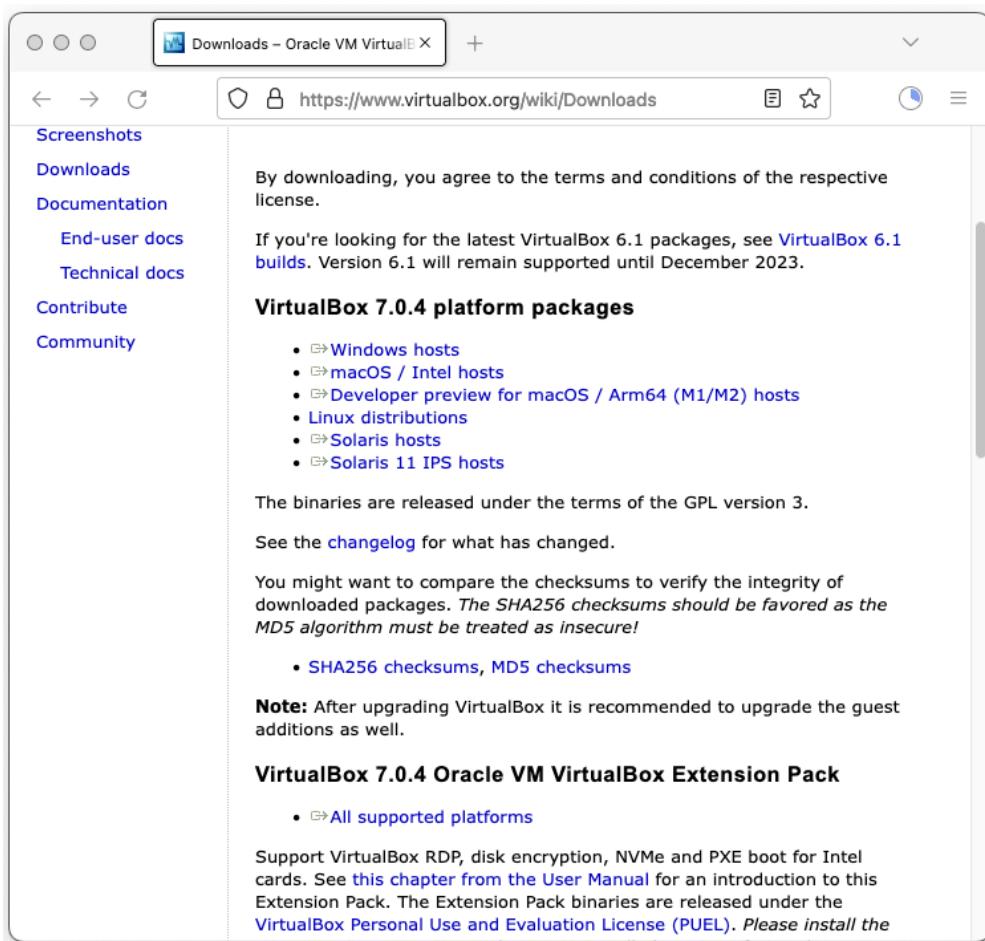


Figure 17.5: The VirtualBox download page

After clicking a download link (Windows or macOS), the VirtualBox installer should start to download:

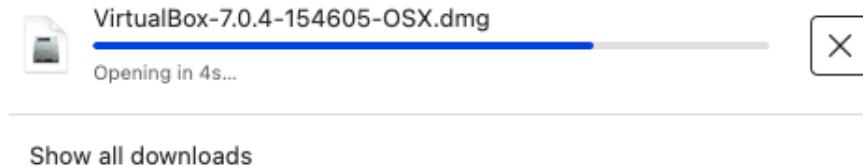


Figure 17.6: The VirtualBox download prompt in Firefox

3. Once VirtualBox has downloaded, you should install it in the usual way that software is installed for your operating system.

Note

On Windows, you may see a warning during installation that says **Missing Dependencies Python Core / win32api**. If this happens, stop the installation and install the pywin32 Python module with the Command Prompt, using the `pip install pywin32` command.

macOS security settings

In macOS, the system security settings may prevent VirtualBox from installing. You will get a message in the installer that says **The Installation Failed**, and macOS will prompt with an alert that says **System Extension Blocked**. Click the **Open Security Preferences** button on the alert and quit the installer.

In the **System Settings** app, the **Security & Privacy** pane should be open. Click the lock button in the bottom-left corner of the window, and then enter your password. Next, click the **Allow** button next to the **System software from developer “Oracle America, Inc.” was blocked from loading** message. After allowing the extension, rerun the installer, and it should complete successfully.

4. Launch VirtualBox. Once again, on macOS, you may have to confirm some extra security settings (Linux/Windows users, continue to *step 5*). You may be presented with an **Accessibility Access (Events)** dialog box (*Figure 17.7*). This may not come up until you try to start the virtual machine, so refer back to these steps if this occurs then:

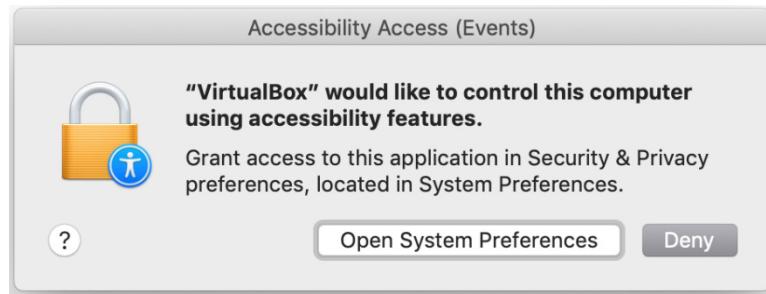


Figure 17.7: The Accessibility Access (Events) dialog

5. Click **Open System Preferences**. The system preferences will open on the **Security & Privacy** pane. Click the lock icon in the bottom-left corner (*Figure 17.8*). You will be prompted to enter your password:

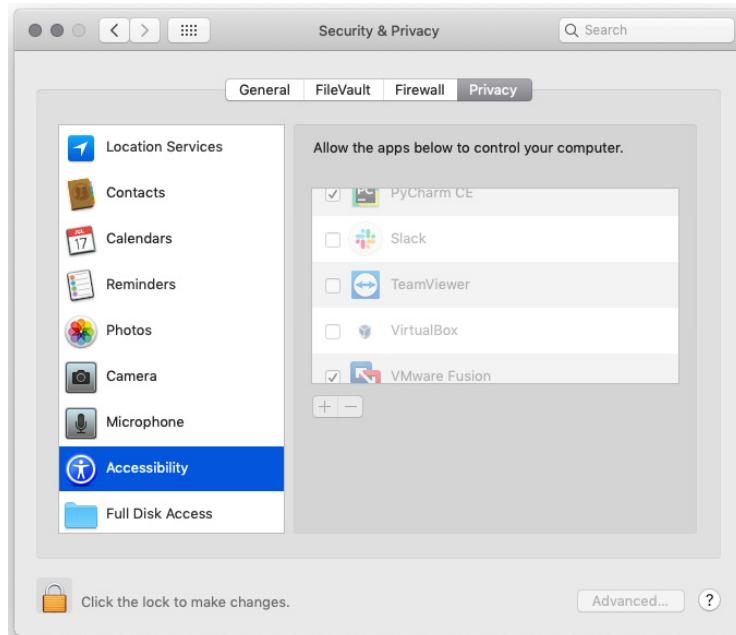


Figure 17.8: The Security & Privacy pane (macOS)

6. Select **Accessibility** on the left pane if it isn't already selected. Then, check **VirtualBox** on the right-hand side under the **Allow the apps below to control your computer.** pane (*Figure 17.9*):

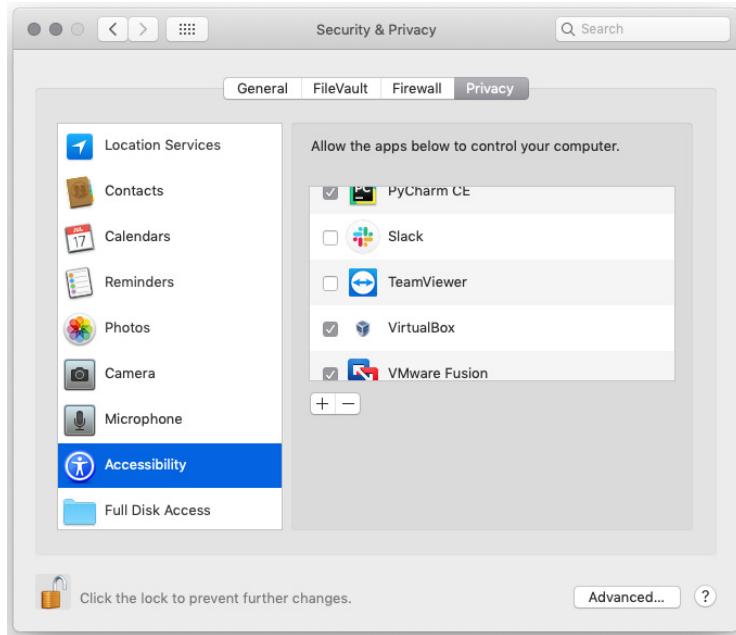


Figure 17.9: VirtualBox enabled in the Accessibility pane

You can now quit the system preferences.

7. Once VirtualBox is up and running, you'll see the **Oracle VM VirtualBox Manager** window:



Figure 17.10: Oracle VM VirtualBox Manager

- Click the **New** button at the top of the window to create a new virtual machine.

You will be stepped through the setup wizard. The first step is the **Name** and operating system settings. Give the virtual machine a name, select where to save it on your computer, and tell VirtualBox what operating system you are installing:

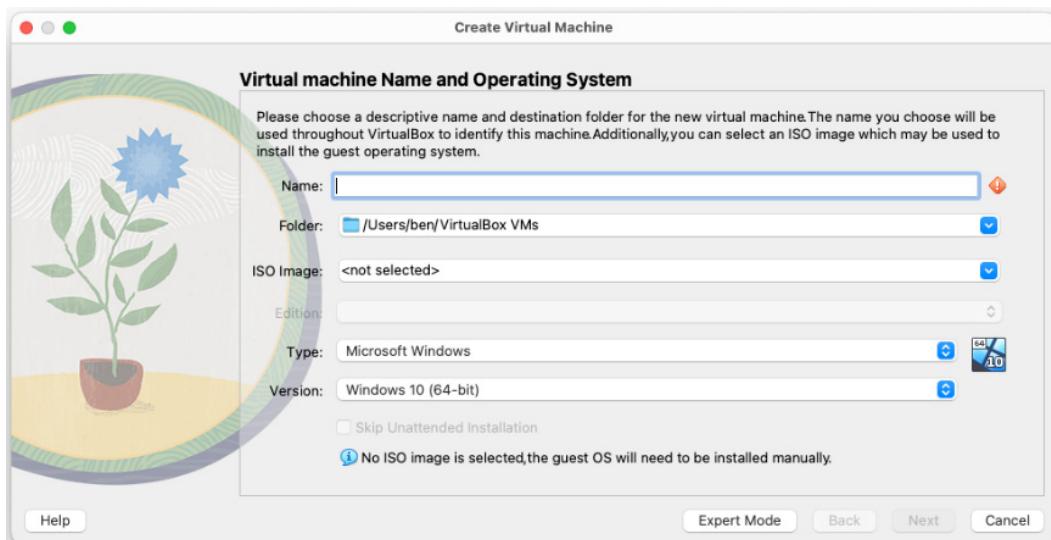


Figure 17.11: The VirtualBox name and operating system setup

The name is used to identify the machine as yours, but it's not very important what you pick. Just enter something such as Bookr.

Folder is where VirtualBox stores the virtual machine's data. The default location is usually fine. The virtual machine's virtual hard drive can take up a few GB of space, so you might choose to store it on another disk if you require more space and have one available.

At this stage, you can specify the ISO image for the virtual machine. This is the file named something such as `ubuntu-22.04.1-live-server-amd64.iso` that you downloaded in *step 1*. This is the equivalent of inserting a CD into a CD drive or attaching a USB flash disk to a computer.



Figure 17.12: VirtualBox ISO image selection

Now that you have specified an ISO image, the **Type** and **Version** fields will be automatically set to **Linux** and **Ubuntu (64-bit)** respectively.

9. Select **Linux** for **Type** and **Ubuntu (64-bit)** for **Version**, and then click **Continue**.

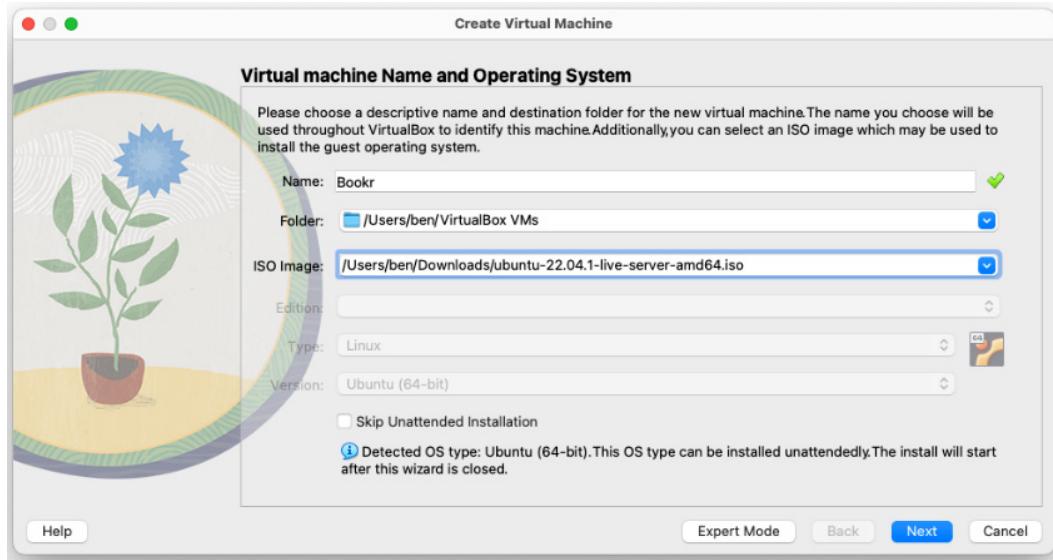


Figure 17.13: The operating system type and version populated

10. Click **Next** to go to the following frame of the wizard.
11. The next frame of the wizard is **Unattended Guest OS Install Setup**. Here, details are entered so that the remainder of the installation can happen in a largely non-interactive fashion.

The **Additional Options** group is for Microsoft product licensing and can be ignored with the current Linux installation.

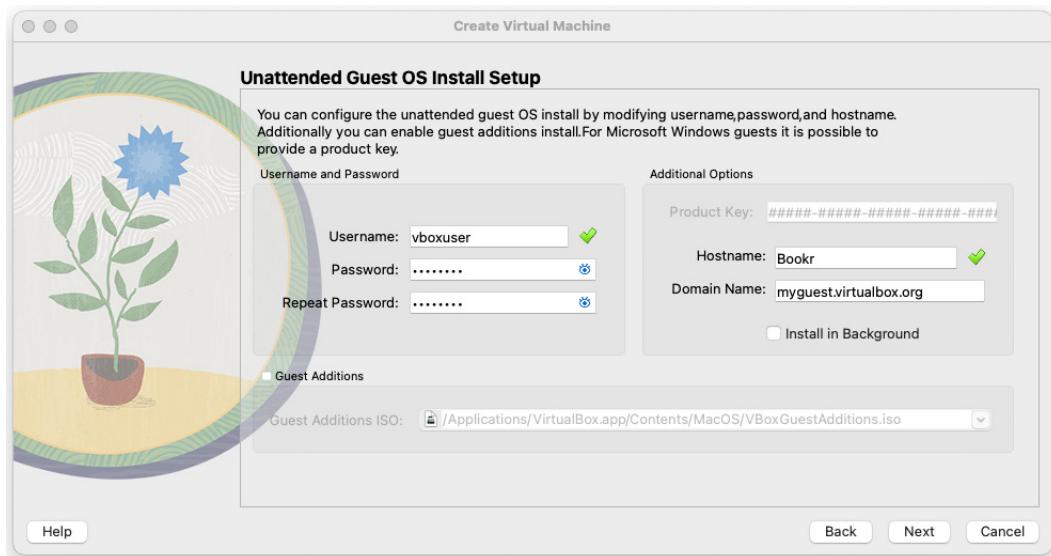


Figure 17.14: VirtualBox – Unattended Guest OS Install Setup

The **Username** and **Password** group is where you specify the details of the non-root account that will be created during setup.

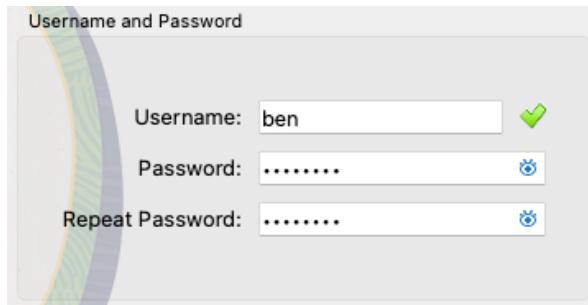


Figure 17.15: VirtualBox – Username and Password

12. Click **Next** for the following frame of the wizard.
13. The next part of the wizard is the memory size, where you set the amount of memory to allocate to your virtual machine:

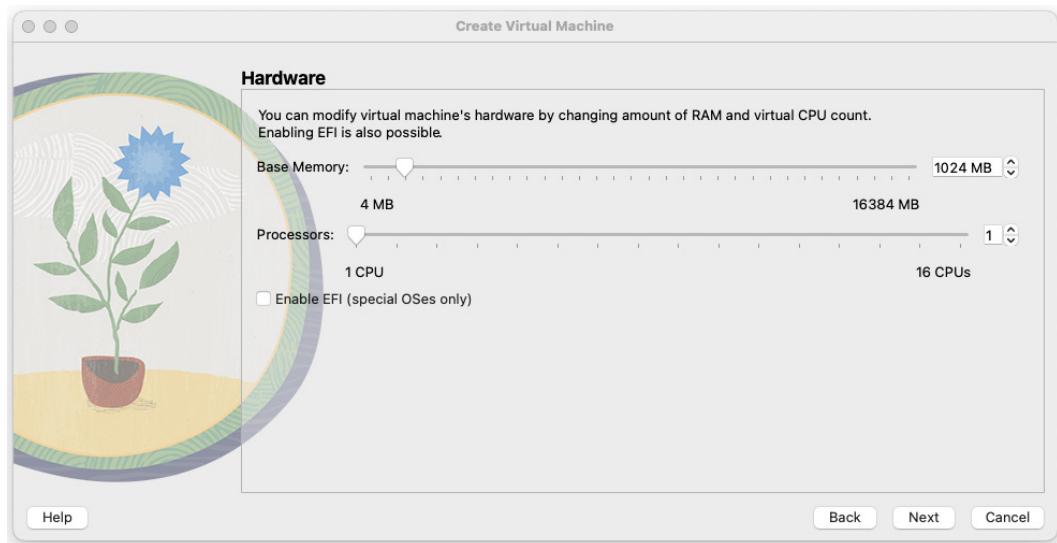


Figure 17.16: Memory size

The more you allocate to it, the less is available to your host computer, and it may cause other applications you're running to slow down. However, if you don't allocate enough, your virtual machine might slow down or the applications it's running may crash.

This default base memory is 2,048 MB (2 GB) and the default number of processors is 1 CPU. The recommended minimum for Ubuntu is 1,024 MB (1 GB), so it should stay at this setting or higher at least during installation. After installation, this can be decreased if it affects the host system's performance.

Note

Choosing the amount of memory for a server can be complicated and depends on how many users you have and how complicated your site and database are. With a real hosted virtual server, you will need to perform testing and scale up or down, based on load and cost requirements.

Click **Next**.

14. Next, set up the virtual hard disk in the **Hard disk** settings. The default option is **25.00 GB**, although we could make do with **10.00 GB** for a development server or if there is no intent to store a large amount of data, such as a big database, multimedia, and log files. Just make sure **Create a Virtual Hard Disk Now** is selected:

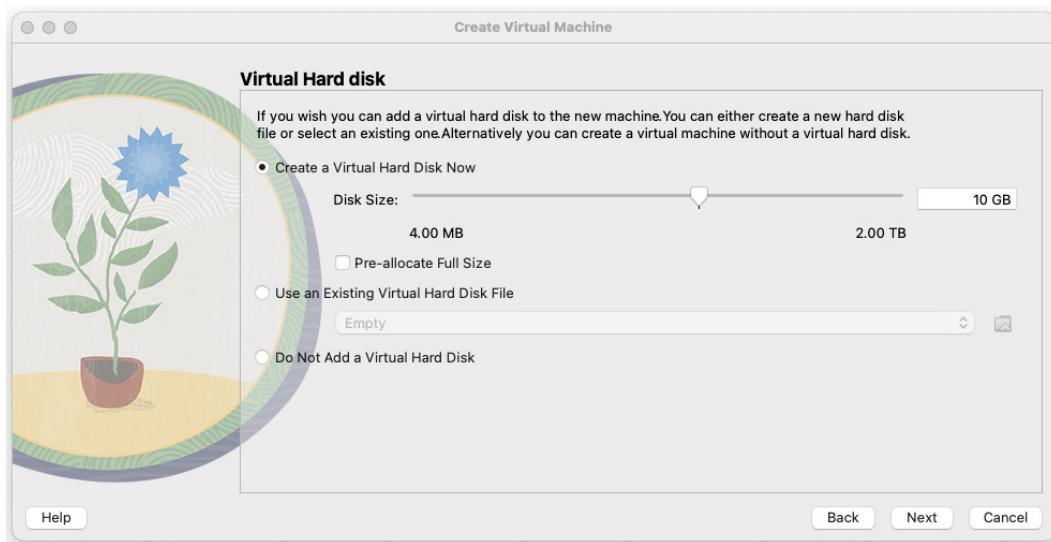


Figure 17.17: The hard disk settings

Then, click **Next**.

15. You will be presented with a summary of your virtual machine configuration, and you can click **Finish** to complete the wizard (*Figure 17.18*):

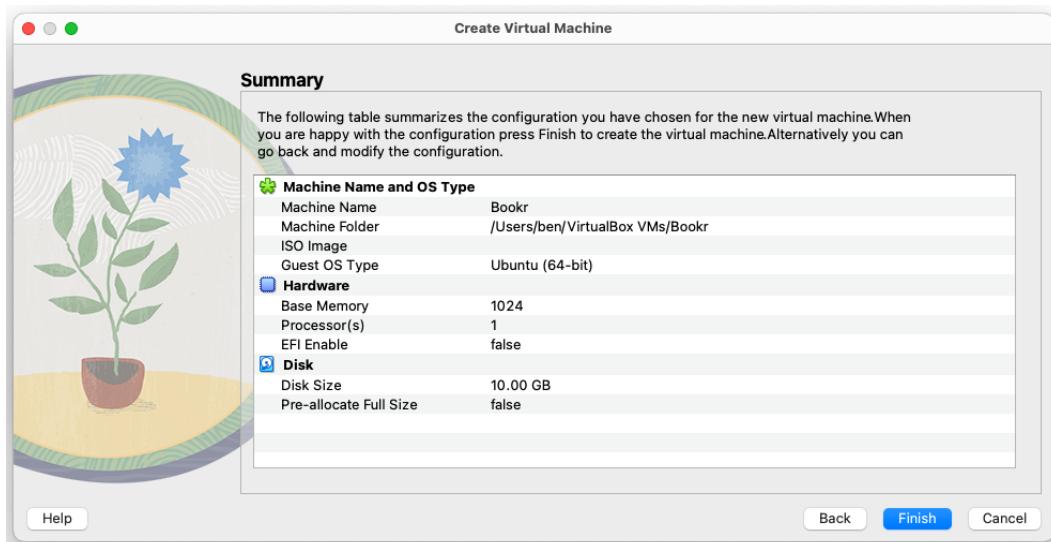


Figure 17.18: The VirtualBox configuration summary

The **Oracle VM VirtualBox Manager** window should now show all the information about the virtual machine we just created:

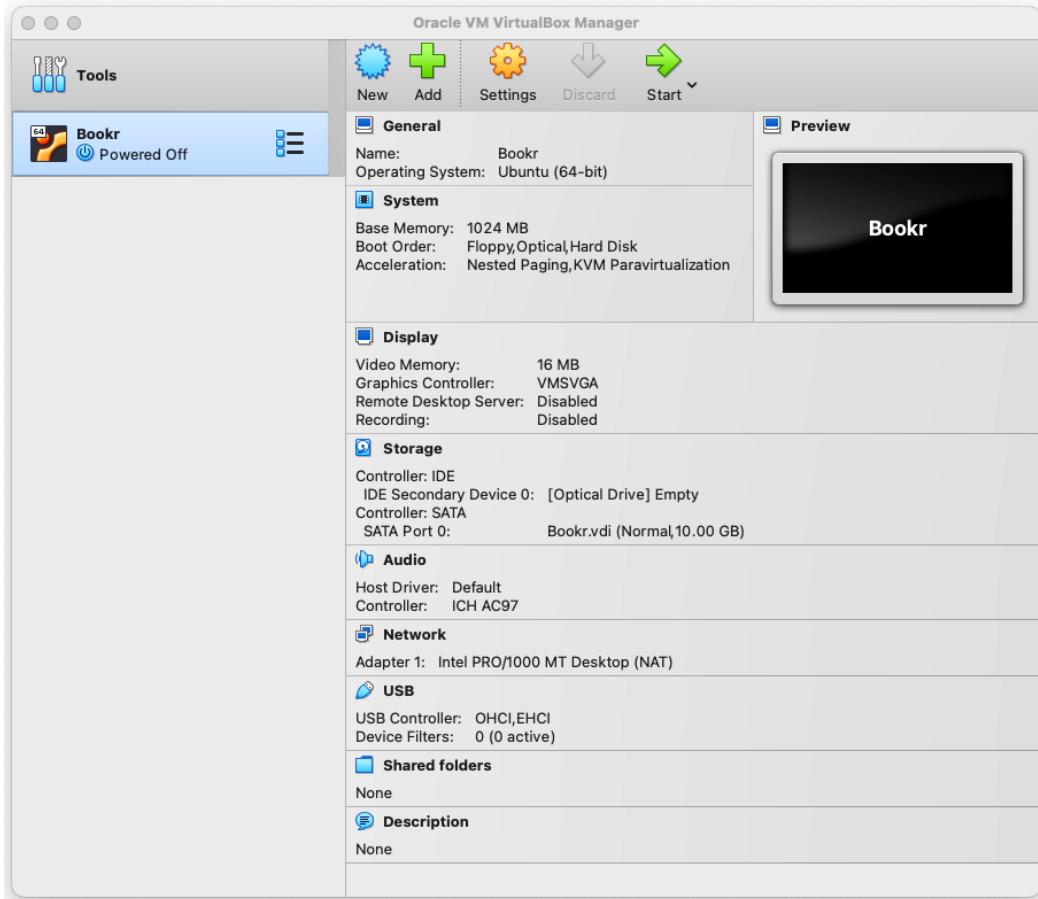


Figure 17.19: The Bookr virtual machine summary

Note

If you're running VirtualBox on macOS with a Retina display, the virtual machine display will be very small. Open the display preferences for the virtual machine by clicking the **Display** section header. Then, set the scale factor to 200% and click **OK** (refer to *Figure 17.20*).

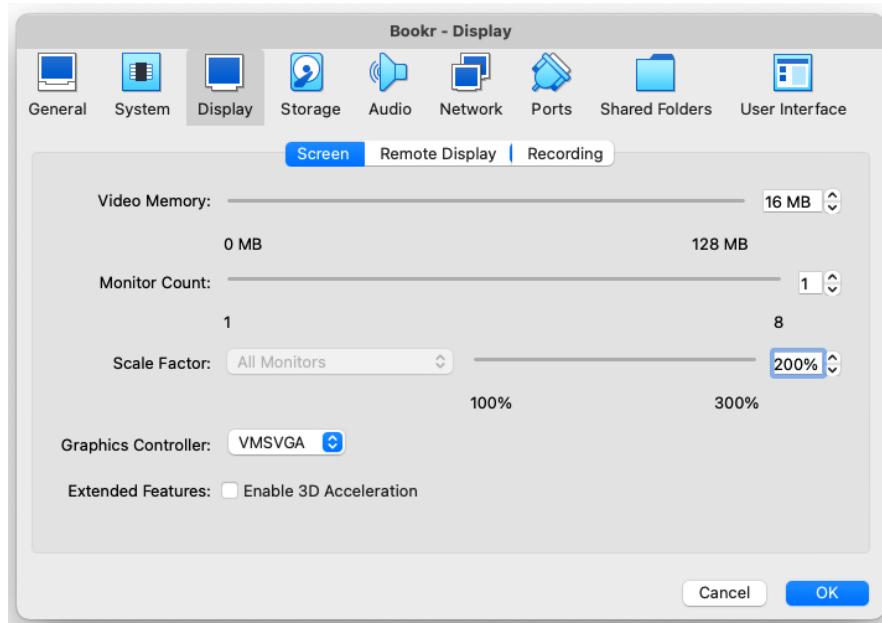


Figure 17.20: The Display scale settings

We're now ready to start the virtual machine, which will begin the Linux installation process. We will continue that in the next exercise.

In this exercise, we downloaded VirtualBox and an Ubuntu Linux ISO in preparation to set up your virtual server. We created a new virtual machine and configured it with enough RAM and hard drive space to run a basic Ubuntu server. It's now ready to boot up and begin the installation.

In the next exercise, we will start up the virtual machine and begin the Ubuntu Linux installation. You can follow these instructions if you have a virtual machine ready to go in VirtualBox or some other virtualization software (such as VMware or Parallels). Alternatively, if you already have a virtual server (or a real Linux server) set up with a cloud provider, you can skip this exercise and use that instead, and move on to *Exercise 17.04 – package update and installation*.

Exercise 17.02 – Ubuntu Linux installation

This exercise continues from *Exercise 17.01 – VirtualBox installation and virtual machine creation*. In the previous exercise, you created a virtual machine (the virtual hardware, so to speak), and in this exercise, you will install the operating system (Ubuntu Linux) on it. You should mostly be able to follow the installation instructions even if you use different virtualization software. If you have an account with a cloud provider that lets you start up a prebuilt Ubuntu Linux server, you can use that instead and skip this exercise:

Note

Throughout this exercise, you will be prompted to press *Enter*. On macOS, you can press *Return* instead.

1. Launch VirtualBox if it's not already running. You should see the **Oracle VM VirtualBox Manager** window.
2. Click the green **Start** button at the top of the **Oracle VM VirtualBox Manager** window. It will boot from the ISO you attached in *Exercise 17.01 – VirtualBox installation and virtual machine creation*. A new window will open, showing what the virtual machine's screen displays.



Figure 17.21: The VirtualBox splash screen

There are two notifications that appear as modal boxes in the right-hand panel of the display. These are to inform you that mouse gestures will be captured by the guest operating system when you hover over the virtual machine window, and that keystrokes will be directed to the guest operating system when this window has focus. You may initially find this user interface behavior disconcerting. Your mouse pointer will disappear from view, as the server installation is console-only and doesn't run a Windows manager containing mouse information; however, if you have "Mouse Integration" configured, your mouse will return to the host OS when it is shifted out of the guest OS window.

Similarly, you can't switch to other programs using normal keystrokes, such as *Alt + Tab* or *Command + Tab*, as these keystrokes are passed directly by the guest operating system (Ubuntu Linux) and are not interpreted by the host operating system (Windows or macOS).

Fortunately, VirtualBox defines a host key to return control of the mouse and keyboard to the host operating system. In Windows, the host key is defined as the right control key, and on macOS, it is defined as the left command key. The host key is displayed in the bottom right of the virtual machine window.

3. Then, you will see an options menu titled **GNU GRUB (GRand Unified Bootloader)**.

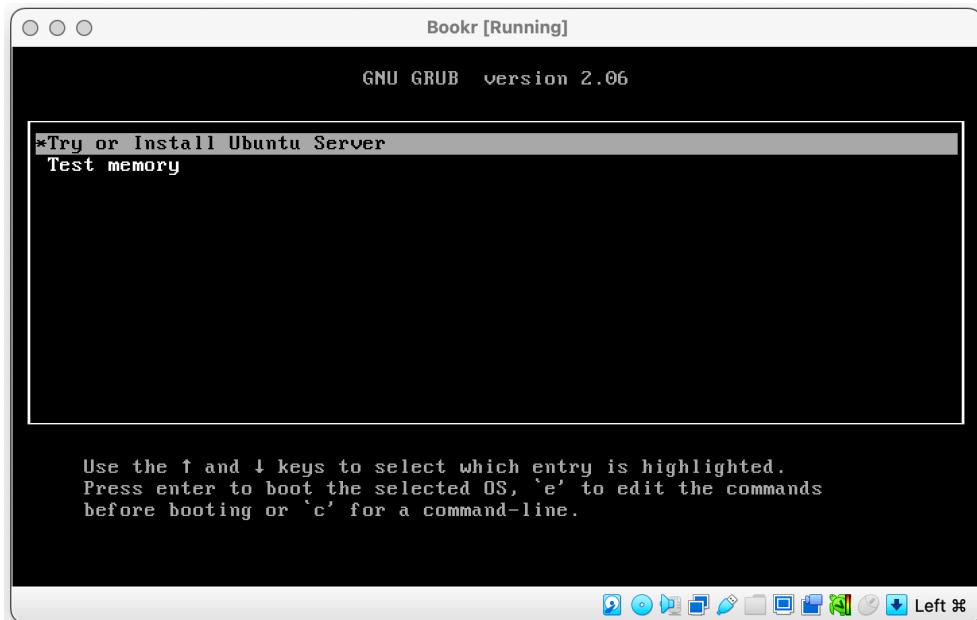


Figure 17.22: The GNU GRUB

This bootload screen will always appear before your Ubuntu VM begins running. If you are unable to gain root access to your Ubuntu VM because of lost passwords or a severe configuration error, you can edit the options in this bootloader menu to reset passwords or repair your system.

Clicking on the window and VirtualBox will capture your mouse and keyboard input. You might get a message like the one shown in *Figure 17.23*:

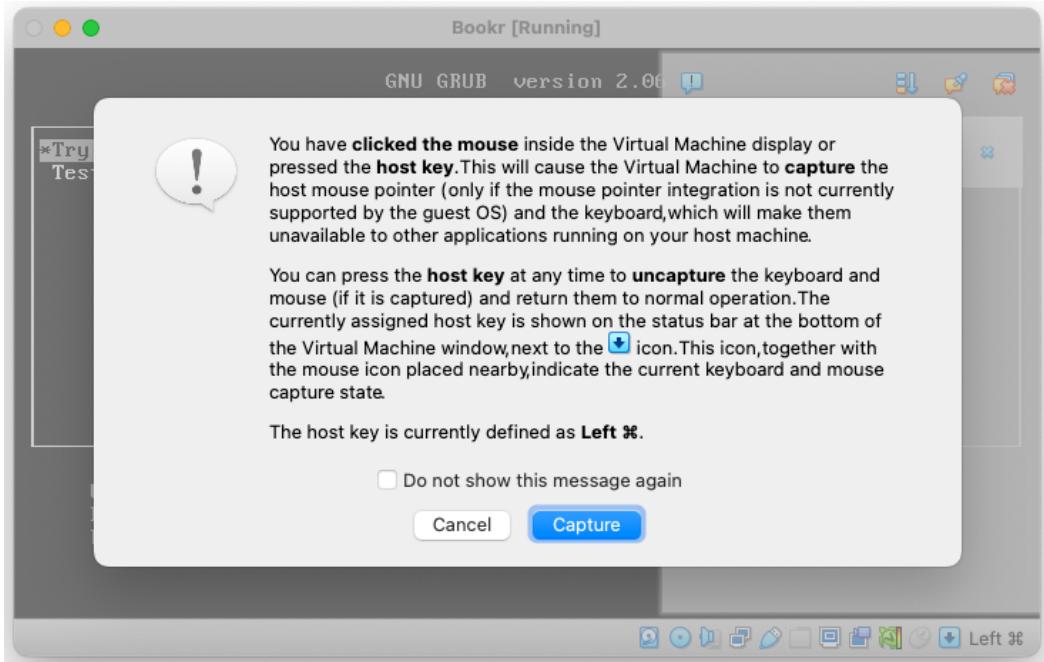


Figure 17.23: The message on the screen after clicking the window

You will have to press the left *Command* key (macOS) or the right *Windows* key (Windows) to release the capture. Hitting *Enter* when **Try or Install Ubuntu Server** is selected will commence the installation process.

4. Ubuntu will start booting up on the virtual screen; you will see some log messages, and within a minute or 2, you will get the language selection screen:

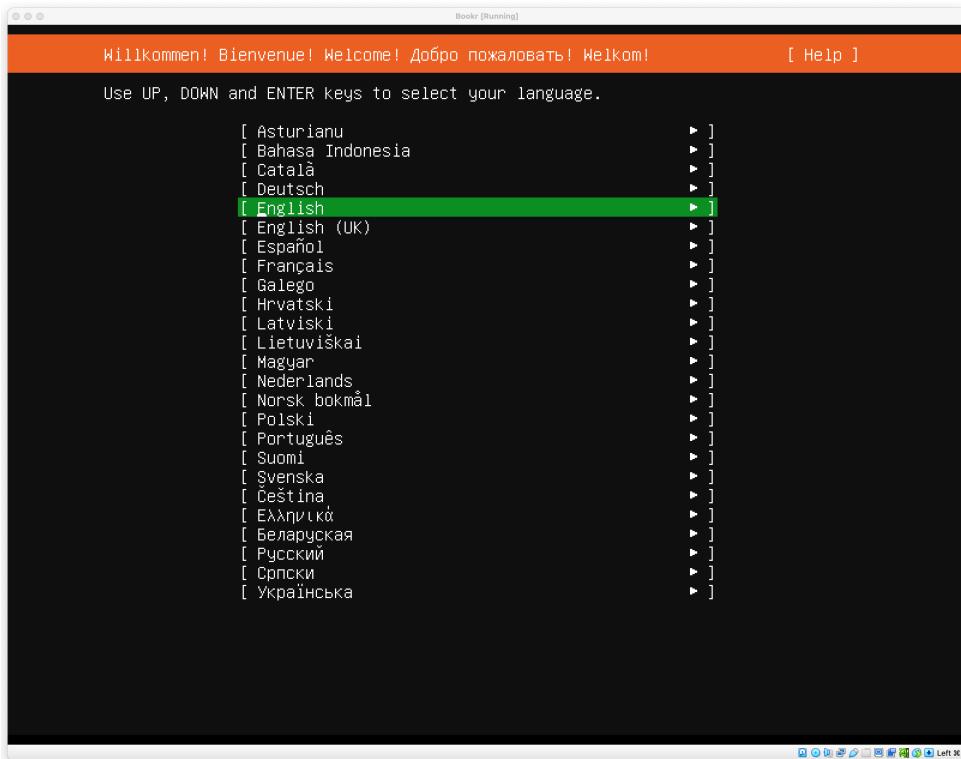


Figure 17.24: The Ubuntu language selection screen

5. Use the up and down cursor keys to select your preferred language, and then press *Enter* to continue.

In the next step, you will be asked whether you want to update the installer before continuing:

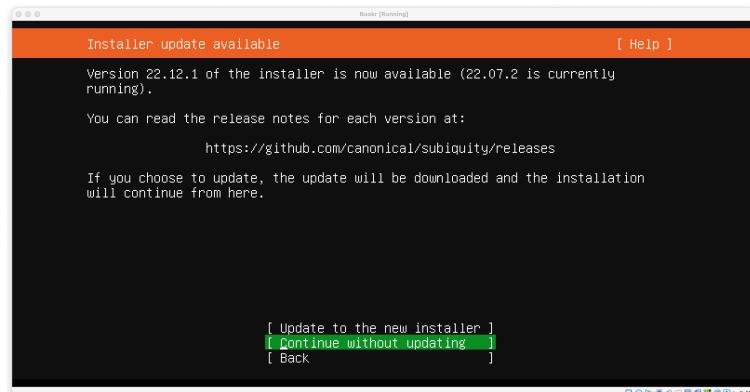


Figure 17.25: The installer update selection

We don't need to do that, so use the cursor keys to select **Continue without updating**, and then press *Enter*.

The installer will automatically select the keyboard layout it thinks you are using:

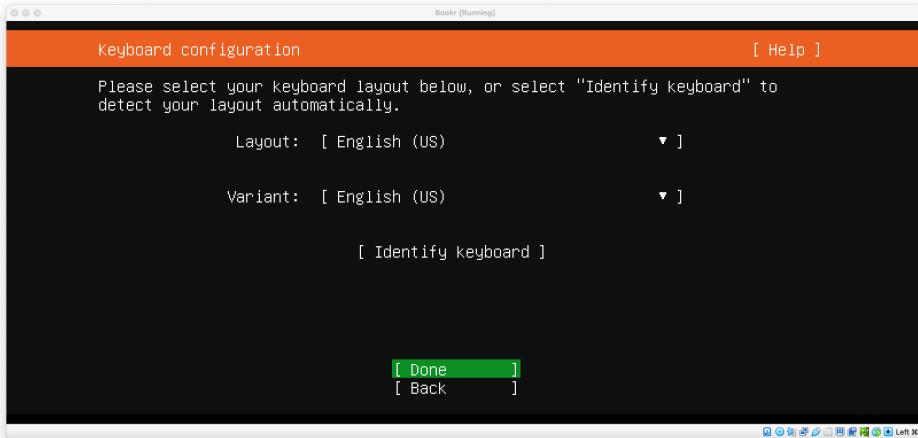


Figure 17.26: The keyboard identification screen

If it doesn't seem to be correct, choose **Identify keyboard**, press *Enter*, and then follow the steps to select a keyboard layout (this involves answering a few questions about your keyboard). When you're happy with the keyboard selection, select **Done** and press *Enter*.

6. Choose the type of installation:

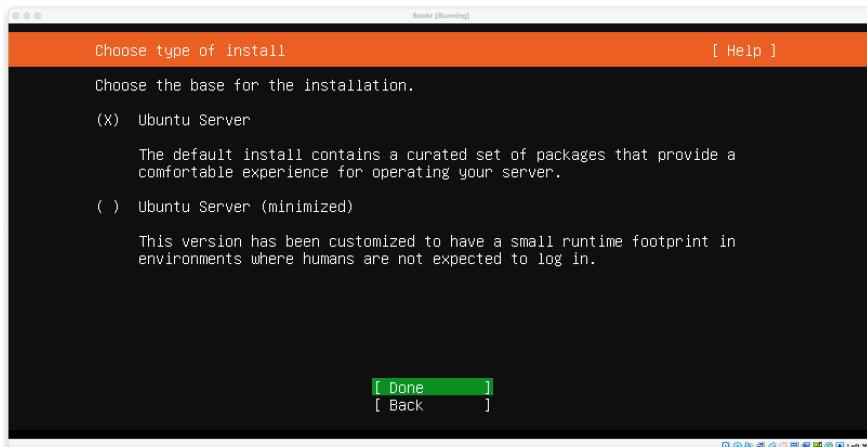


Figure 17.27: Choosing the type of installation

7. You'll need to select your network interface in the next step:

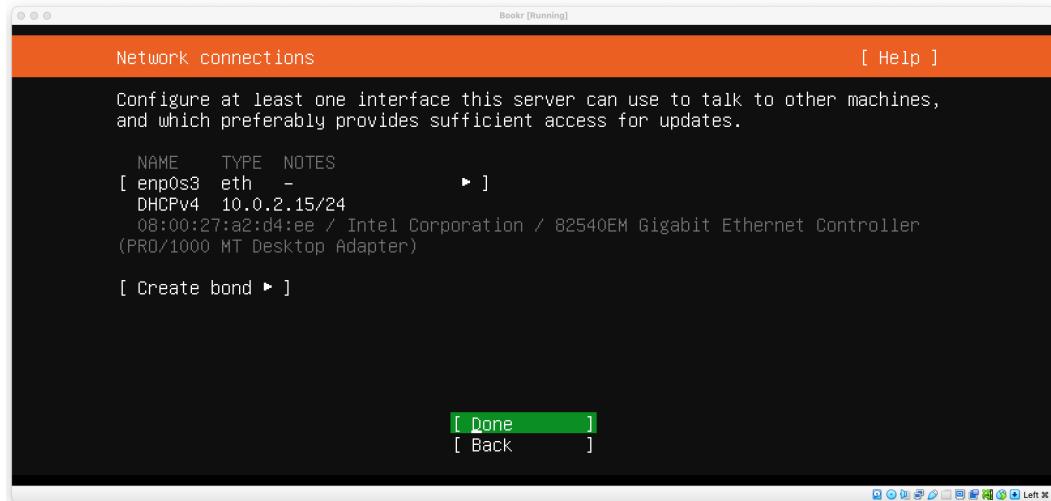


Figure 17.28: The network connection setup

Since VirtualBox provides standard virtual hardware with only one network interface, just select **Done** and press *Enter*.

8. On the next screen (*Figure 17.29*), you can enter a proxy server:

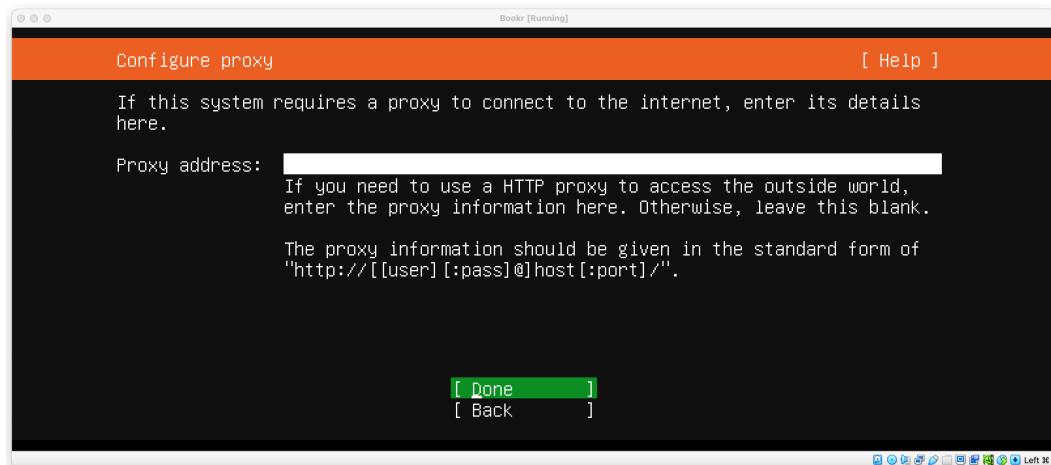


Figure 17.29: The proxy server configuration

You should type in a proxy server if you need to use one to access the internet. If not, just select **Done** and press *Enter*.

9. Next, the Ubuntu installer will automatically pick the mirror server to download packages from; this is based on your location:

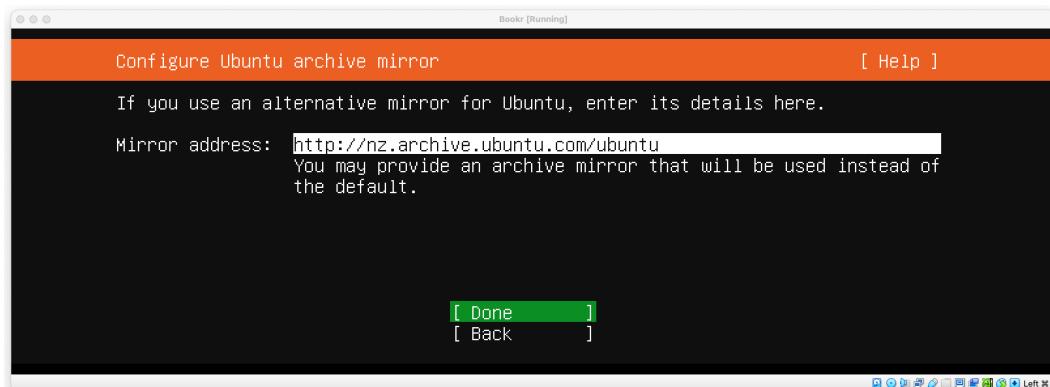


Figure 17.30: The mirror server selection (auto-selected)

You can leave this as the default (it is selected automatically based on your region), and then select **Done** and press *Enter*.

The next step is to partition the virtual hard drive and format it. The installer guides us through this process, and since we're using the entire virtual disk, it's very simple:

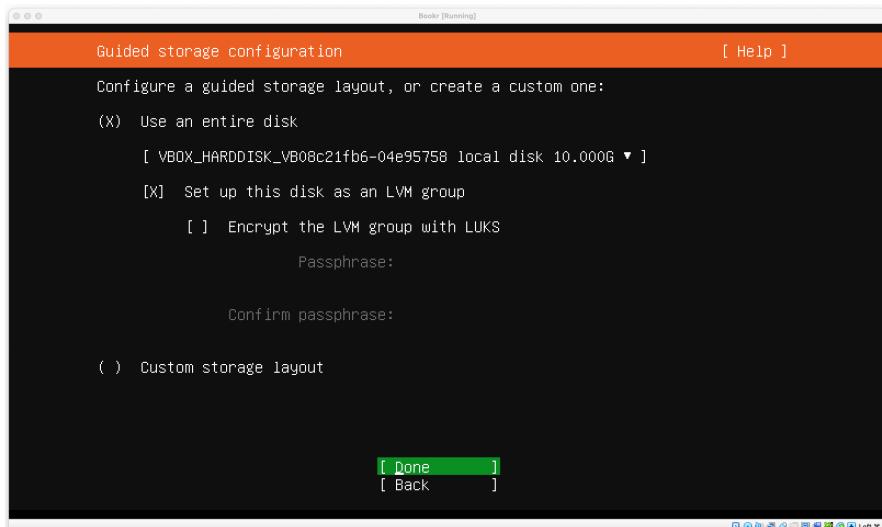


Figure 17.31: The filesystem setup

10. Select **Use an entire disk**, and then press *Enter*.
11. Then, select the hard disk to use. Since there's only one, select it from the list (it will start with **VBOX_** – see *Figure 17.32*), and then press *Enter*:

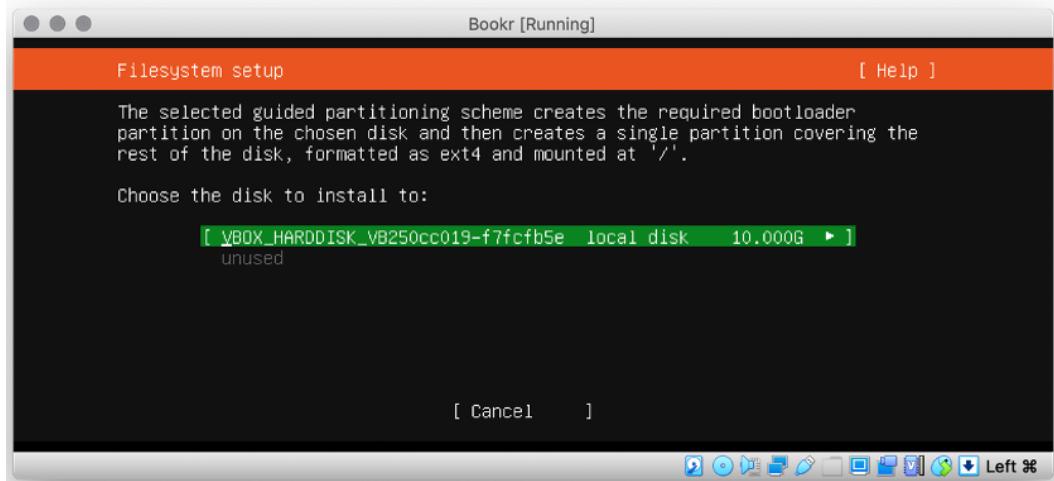


Figure 17.32: Hard disk selection (no corresponding)

It should be noted here that this applies to your virtual hard disk. No changes will be made to the formatting of the host computer's disk. Even if something goes wrong during the installation, you can just delete the virtual disk and start again.

12. On the next screen (*Figure 17.33*), the installer automatically generates a partition scheme. Just select **Done** and press *Enter*:

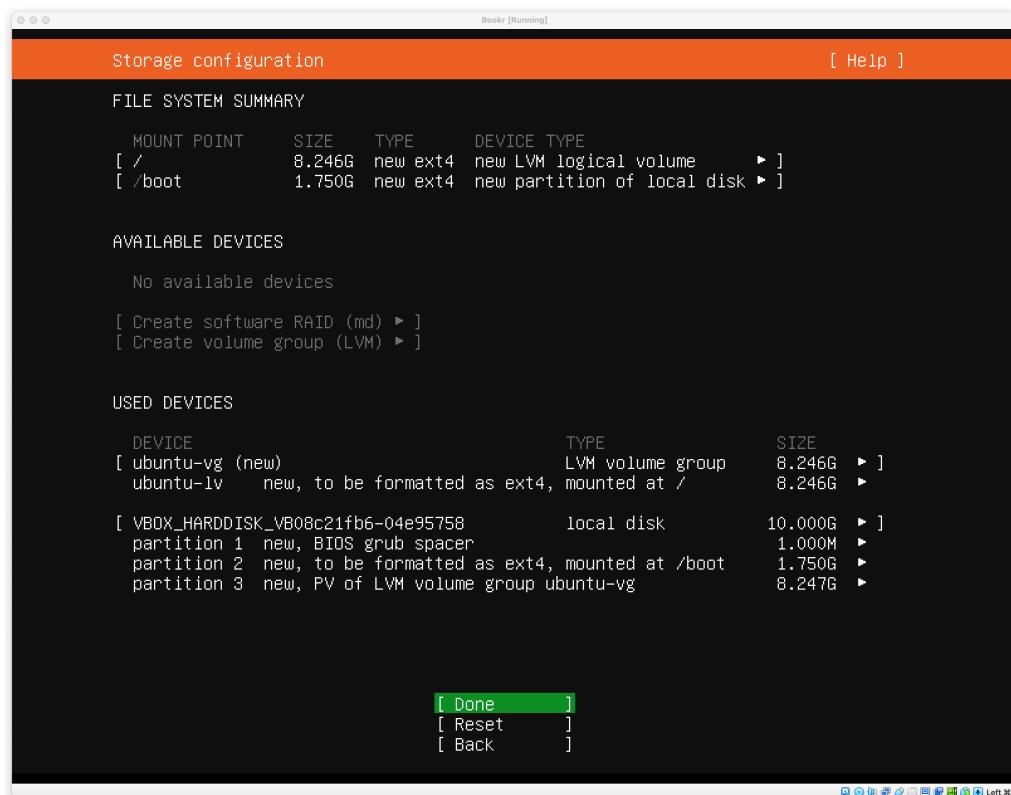


Figure 17.33: The partition entry

13. You will be asked to confirm that you want to format the disk. Use the arrow keys to select **Continue**, and then press *Enter*:

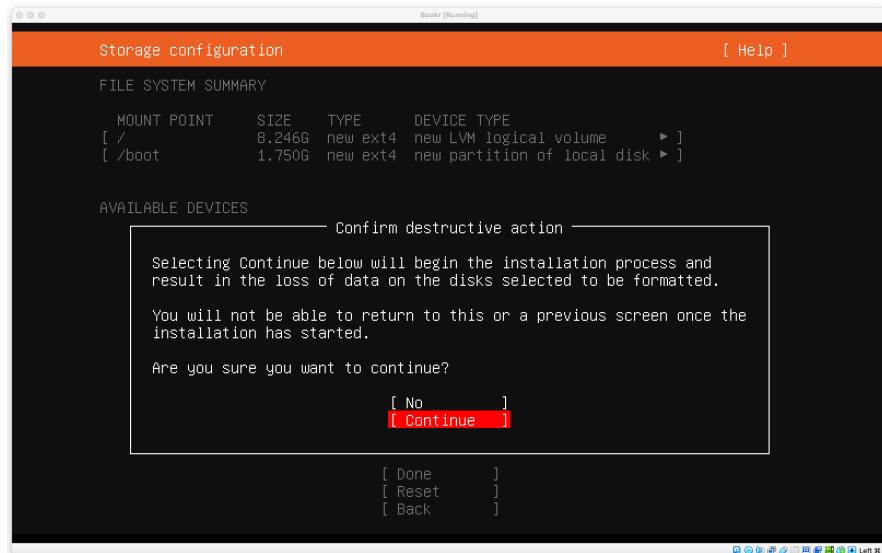


Figure 17.34: Confirming the disk formatting

14. After the formatting is complete, on the next screen, you will need to enter your details. After typing in each field, use *Tab* to move to the next field. See *Figure 17.35* for an example:

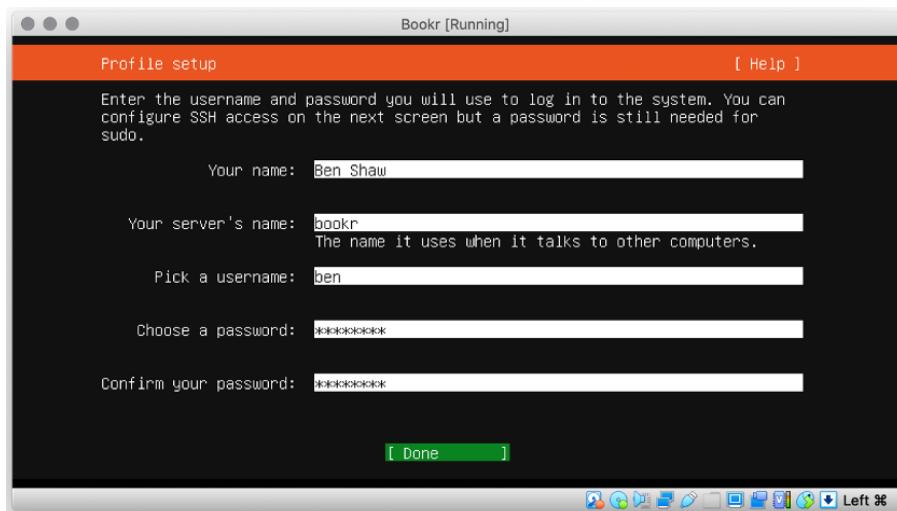


Figure 17.35: The profile setup

Complete the fields as follows:

- Your name.
- For Your server's name, enter bookr.
- For your username, you can choose your first name or any short username you like. This is the name of your personal user – we will create a specific Bookr user later on, after the server is set up.
- Enter a password for your user account, and do the same in the **Confirm your password** field.
- Press *Tab* again to select **Done**, and then press *Enter*.

On the next screen, we want to enable SSH access on the server so that we can connect to it using SSH and SFTP:

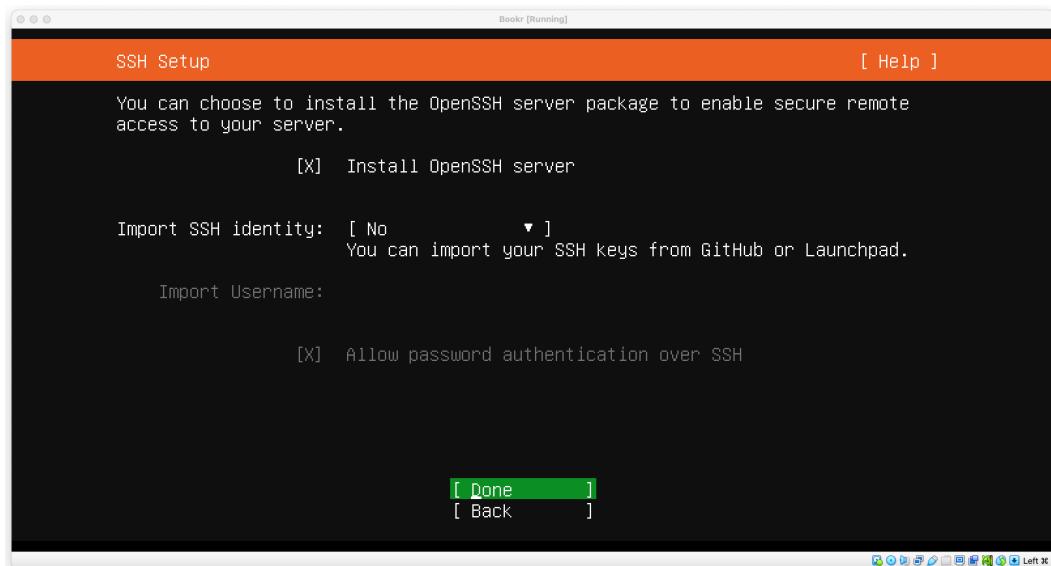


Figure 17.36: The SSH setup

15. Press the spacebar to check the **Install OpenSSH server** option. Then, move the cursor down, select **Done**, and press *Enter*.

The next screen allows you to select packages to pre-install:

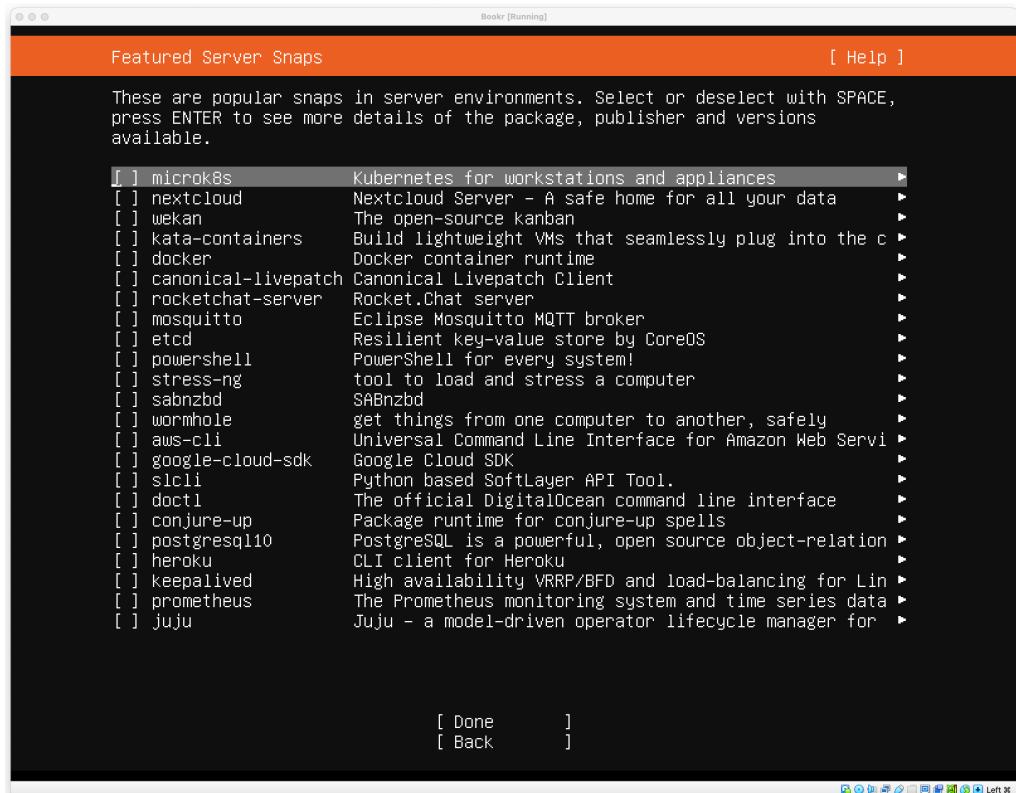
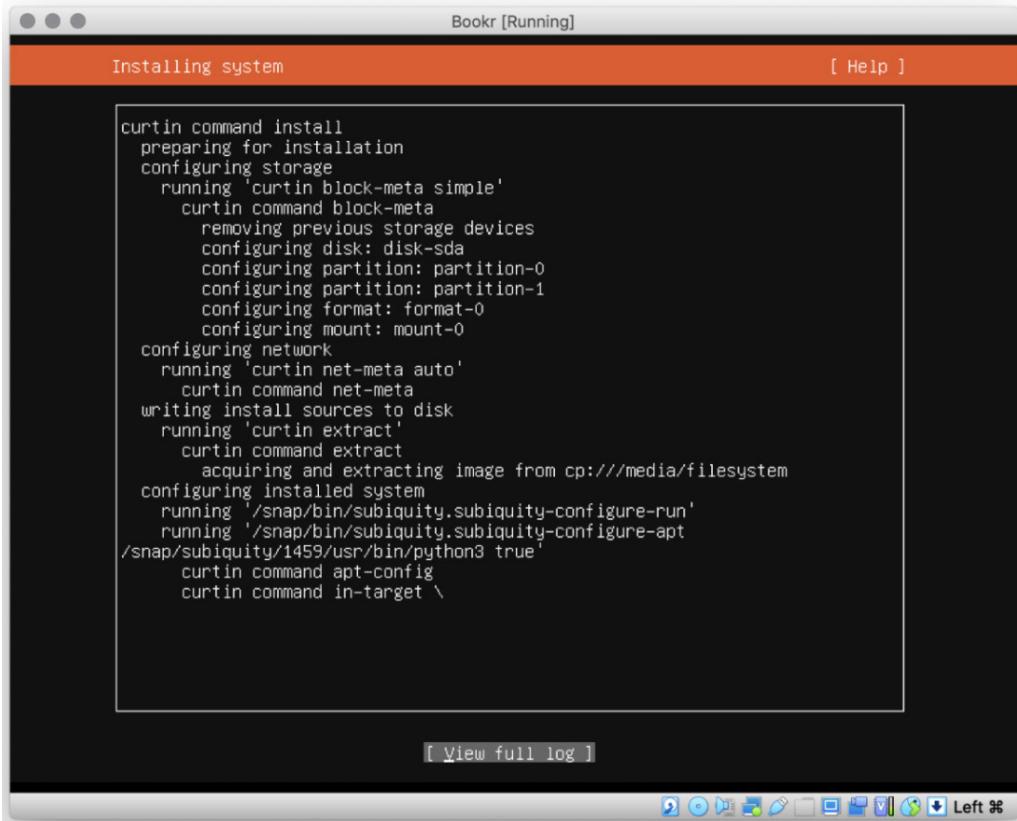


Figure 17.37: The package selection

16. Note **postgresql10** in this list. While we could install PostgreSQL at this step, for consistency, we will install it later with all the other packages we need. Since we don't want to install anything now, move down to select **Done** and press *Enter*.

17. The next screen will show a log of the activity, and the installer will perform the base installation and then automatically install the latest security updates:



The screenshot shows a terminal window titled "Bookr [Running]" with the title bar color set to orange. The window contains a black background with white text displaying a log of system installation commands. At the bottom of the window, there is a small button labeled "[View full log]". The window has a standard window frame with close, minimize, and maximize buttons at the top left, and a toolbar with various icons at the bottom right.

```
curtin command install
  preparing for installation
  configuring storage
    running 'curtin block-meta simple'
      curtin command block-meta
        removing previous storage devices
        configuring disk: disk-sda
        configuring partition: partition-0
        configuring partition: partition-1
        configuring format: format-0
        configuring mount: mount-0
  configuring network
    running 'curtin net-meta auto'
      curtin command net-meta
  writing install sources to disk
    running 'curtin extract'
      curtin command extract
        acquiring and extracting image from cp:///media/filesystem
  configuring installed system
    running '/snap/bin/subiquity.subiquity-configure-run'
    running '/snap/bin/subiquity.subiquity-configure-apt'
/snap/subiquity/1459/usr/bin/python3 true'
  curtin command apt-config
  curtin command in-target \
```

Figure 17.38: The installation screen

This can take some time, so you will need to just wait for it to finish.

When it's finished, the options at the bottom will change to **View full log** and **Reboot Now**:

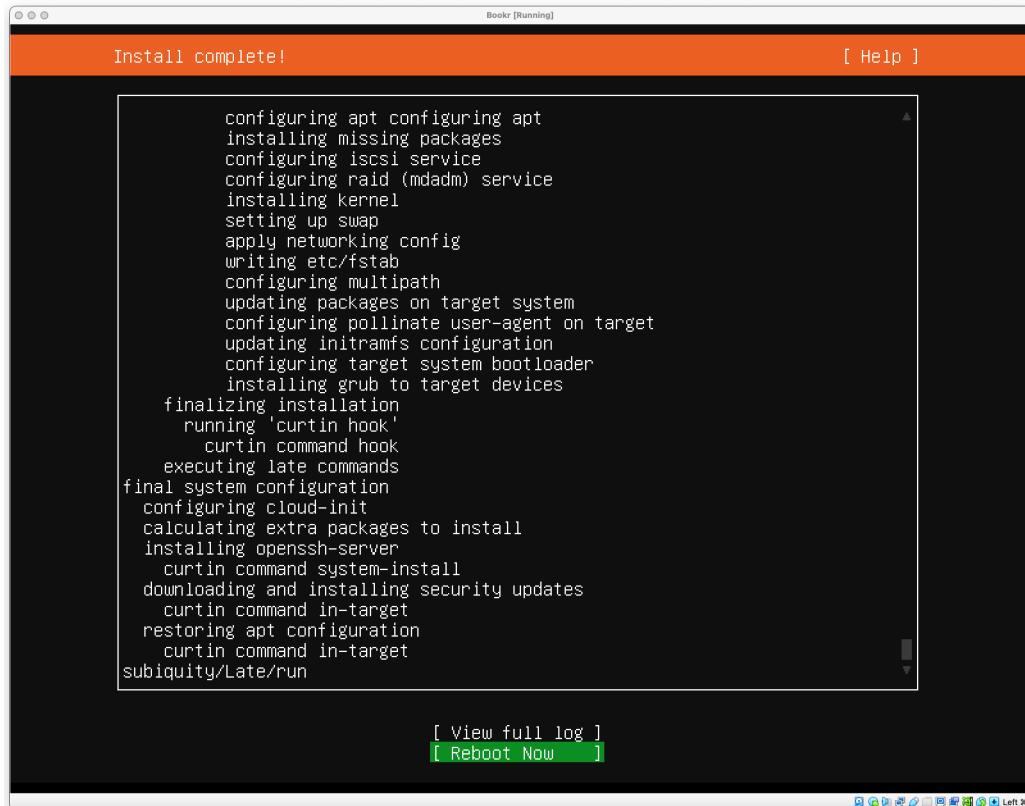


Figure 17.39: Installation complete

18. Select **Reboot Now** and press *Enter*.

The screen will update, start showing more log messages, eventually stop, and then show a **Please remove the installation medium, then press ENTER** message:

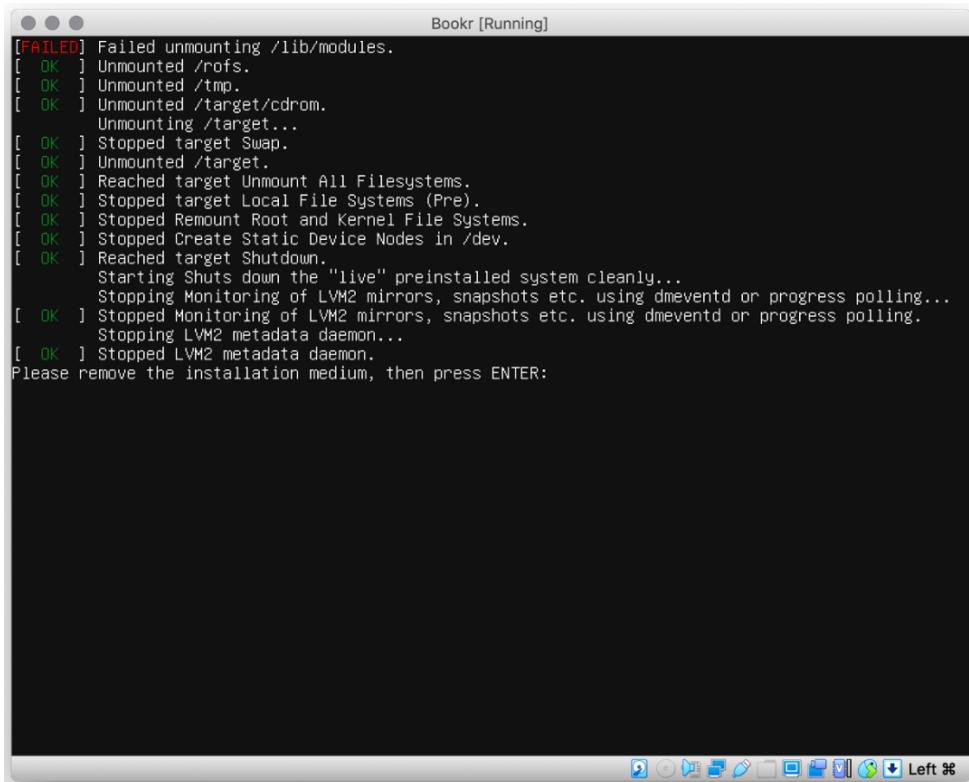


Figure 17.40: The post-installation screen

Note that you might have a **FAILED** message, but that's OK.

19. To “eject” the virtual CD, go back to the **Oracle VM VirtualBox Manager** window and click the **Storage** text header. This will open the **Storage** preferences window, which we saw before, in *step 12*. Once again, click on the CD icon under the **Controller: IDE** heading on the left (this time, it should have the name of the Ubuntu ISO). Then, click on the CD in the right-hand-side **Attributes** section. From the menu that appears (*Figure 17.41*), choose **Remove Disk from Virtual Drive**, and then click **OK** to close and save the storage settings:

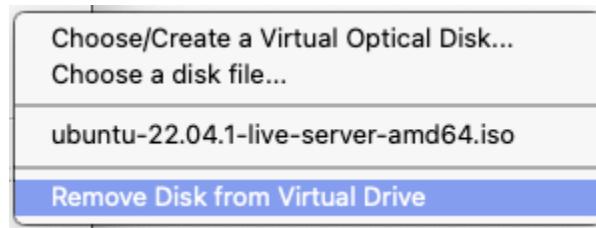


Figure 17.41: Remove Disk from Virtual Drive

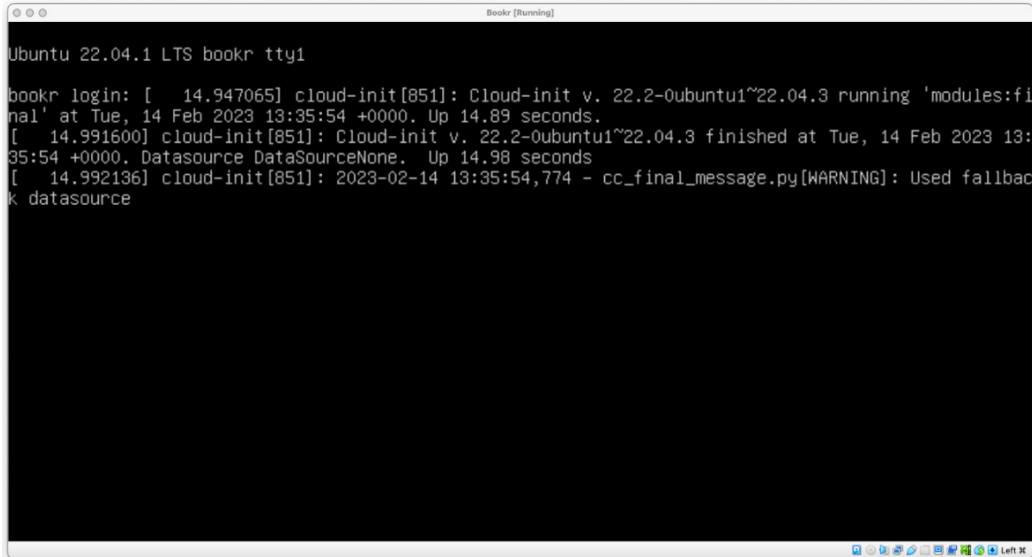
20. Switch back to the Bookr virtual display and press *Enter*. Your virtual machine will reboot, and you should see some log messages as it starts up:

```
[ OK ] Finished File System Check on /dev/d...id/4deb4839-993a-460f-8969-df9a4fb250ab.  
[ OK ] Mounting /boot...  
[ OK ] Mounted /boot.  
[ OK ] Reached target Local File Systems.  
Starting Load AppArmor profiles...  
Starting Set console font and keymap...  
Starting Create final runtime dir for shutdown pivot root...  
Starting Tell Plymouth To Write Out Runtime Data...  
Starting Create Volatile Files and Directories...  
Starting Uncomplicated firewall...  
[ OK ] Finished Create final runtime dir for shutdown pivot root.  
[ OK ] Finished Tell Plymouth To Write Out Runtime Data.  
[ OK ] Finished Uncomplicated firewall.  
[ OK ] Finished Create Volatile Files and Directories.  
Starting Network Time Synchronization...  
Starting Record System Boot/Shutdown in UTMP...  
[ OK ] Finished Record System Boot/Shutdown in UTMP.  
[ OK ] Started Network Time Synchronization.  
[ OK ] Reached target System Time Set.  
[ OK ] Finished Set console font and keymap.  
[ OK ] Finished Load AppArmor profiles.  
Starting Load AppArmor profiles managed internally by snapd...  
Starting Initial cloud-init job (pre-networking)...  
Mounting Arbitrary Executable File Formats File System...  
[ OK ] Mounted Arbitrary Executable File Formats File System.  
[ OK ] Finished Initial cloud-init job (pre-networking).  
[ OK ] Reached target Preparation for Network.  
Starting Network Configuration...  
[ OK ] Started Network Configuration.  
Starting Wait for Network to be Configured...  
Starting Network Name Resolution...  
[ OK ] Started Network Name Resolution.  
[ OK ] Reached target Network.  
[ OK ] Reached target Host and Network Name Lookups.
```

Figure 17.42: Log messages when booting Ubuntu

After a minute or so, you will get the Ubuntu login screen.

Note that if you wait too long before logging in, you might see some `cloud-init` messages printed on the screen, such as those shown in *Figure 17.43*:

A screenshot of a terminal window titled "Bookr [Running]". The window contains the following text:

```
Ubuntu 22.04.1 LTS bookr tty1

bookr login: [ 14.947065] cloud-init[851]: Cloud-init v. 22.2-0ubuntu1~22.04.3 running 'modules:final' at Tue, 14 Feb 2023 13:35:54 +0000. Up 14.89 seconds.
[ 14.991600] cloud-init[851]: Cloud-init v. 22.2-0ubuntu1~22.04.3 finished at Tue, 14 Feb 2023 13:35:54 +0000. Datasource DataSourceNone. Up 14.98 seconds
[ 14.992136] cloud-init[851]: 2023-02-14 13:35:54,774 - cc_final_message.py[WARNING]: Used fallback datasource
```

The terminal window has a dark background and white text. The title bar says "Bookr [Running]". The bottom of the window shows a standard Linux desktop taskbar with icons for various applications.

Figure 17.43: The `cloud-init` log messages

This is OK; just press *Enter* to get the login prompt back.

21. Type in the username you created in *step 14*, and then press *Enter*. Type in your password (again from *step 14*) – note that you won't see any characters appear as you type. Then, press *Enter* to log in. You will see an output of some system information, and then a Command Prompt, such as `yourusername@bookr :~$:`



```
bookr login: ben
Password:
Welcome to Ubuntu 22.04.1 LTS (GNU/Linux 5.15.0-58-generic x86_64)

 * Documentation: https://help.ubuntu.com
 * Management: https://landscape.canonical.com
 * Support: https://ubuntu.com/advantage

 System information as of Sat Jan 14 04:42:04 AM UTC 2023

 System load: 0.123046875 Processes: 104
 Usage of /: 50.4% of 8.02GB Users logged in: 0
 Memory usage: 21% IPv4 address for enp0s3: 10.0.2.15
 Swap usage: 0%

61 updates can be applied immediately.
To see these additional updates run: apt list --upgradable

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

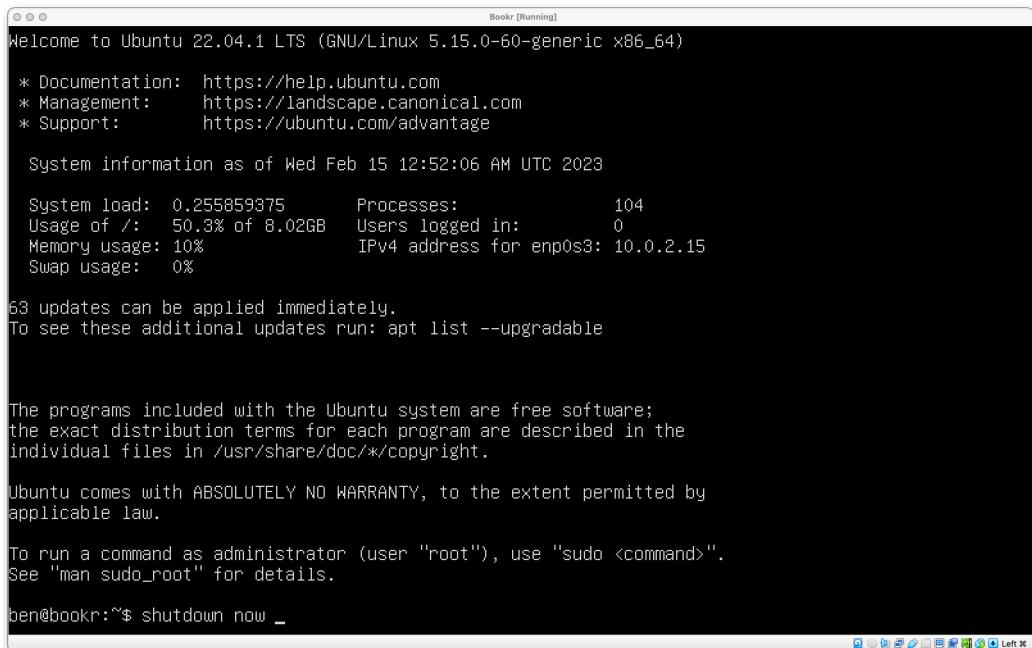
ben@bookr:~$
```

Figure 17.44: Linux Command Prompt

22. You've successfully set up an Ubuntu virtual machine. We'll shut it down from inside the virtual machine, and then come back to it later. To shut the machine down, just type the following:

```
shutdown now
```

The next screenshot shows how your virtual screen should look, with the `shutdown now` command typed in (but before you've pressed *Enter*):



A screenshot of a terminal window titled "Bookr [Running]". The window displays the following text:

```
Welcome to Ubuntu 22.04.1 LTS (GNU/Linux 5.15.0-60-generic x86_64)

* Documentation: https://help.ubuntu.com
* Management: https://landscape.canonical.com
* Support: https://ubuntu.com/advantage

System information as of Wed Feb 15 12:52:06 AM UTC 2023

System load: 0.255859375      Processes:          104
Usage of /:   50.3% of 8.02GB  Users logged in:        0
Memory usage: 10%              IPv4 address for enp0s3: 10.0.2.15
Swap usage:   0%

63 updates can be applied immediately.
To see these additional updates run: apt list --upgradable

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/*copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

ben@bookr:~$ shutdown now _
```

Figure 17.45: The `shutdown now` command entered

23. Then, press *Enter*.

This instructs the machine to power off right now. You could also run `shutdown` with no argument, which would have the machine turn off in 1 minute, or you could specify a later time.

You'll see some shutdown log messages, and then the virtual machine will shut down and its virtual screen will close. In the **Oracle VM VirtualBox Manager** window, you'll see that the Bookr virtual machine is now in the **Powered Off** state:

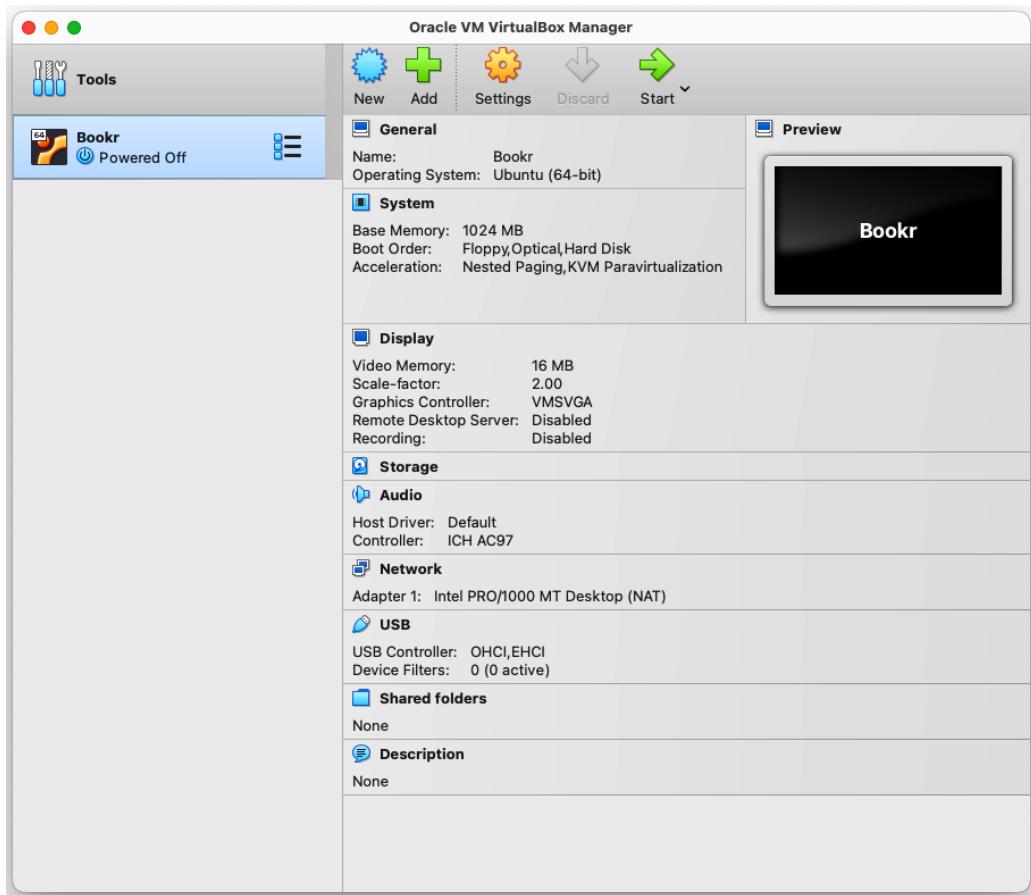


Figure 17.46: The Bookr machine in the Powered Off state

In this exercise, we installed Ubuntu Linux on a VirtualBox virtual machine. This exercise and the previous one were quite long, but you now have an Ubuntu Linux virtual machine that will mimic that from a cloud provider. You will use it in the rest of the exercises and activities.

Now that we have completed an Ubuntu installation, we can familiarise ourselves with running Linux commands. We will use them to carry out some initial tasks for network configuration and software installation.

Working with Linux

In this section, we will take a look at the various Linux commands and understand how to use them.

The sudo command

Later, in *Exercise 17.03 – VirtualBox networking configuration*, we will run a command on the virtual machine using the `sudo` command. First, let's do a brief rundown of some common Linux permissions. We created a Linux user during the virtual machine setup in *Exercise 17.02 – Ubuntu Linux installation*. This is a normal, non-administrative user. There are some commands and files that this user can't access and that are only available to the *superuser*. The superuser has the username `root`, and it's more common to refer to this user as `root` – for example, “*This command needs to be run as root*,” or “*Only the root user can access these files*.” The root user has its own password.

Without going into too much detail, the `sudo` command allows you to execute (do) commands as though you were the superuser (`su`). You authenticate using your own password instead of the root's password. This is more secure than having a root password that is shared among different users, and more fine-grained access can be given to certain users and/or commands.

To run a command as a superuser, you simply prefix it with `sudo`. For example, the `whoami` command simply tells you the name of the current user executing the command (highlighted):

```
ben@bookr:~$ whoami
ben
ben@bookr:~$
```

Next, we can try executing `whoami` using `sudo`. We will be prompted to enter our password. Then, after hitting *Enter*, we'll see it output `root`:

```
ben@bookr:~$ sudo whoami
[sudo] password for ben:
root
ben@bookr:~$
```

After executing a command as `sudo`, our authentication is cached in the current terminal for 5 minutes, so if we execute a command using `sudo` again, we aren't prompted for our password. This applies even if it's a different command.

For example, if we then run `sudo id` (`id` is like `whoami` but shows more information about your user ID and the groups to which you belong), we see `root`'s information without being prompted for a password:

```
ben@bookr:~$ sudo id
uid=0(root) gid=0(root) groups=0(root)
ben@bookr:~$
```

To expire our authentication, we can do any of the following:

- Log out and then log back in
- Wait 5 minutes

- Run sudo -k:

```
ben@bookr:~$ sudo -k
ben@bookr:~$ sudo whoami
[sudo] password for ben:
root
ben@bookr:~$
```

That's all you need to know about sudo for now. We will use it in the next exercise to shut down the virtual machine.

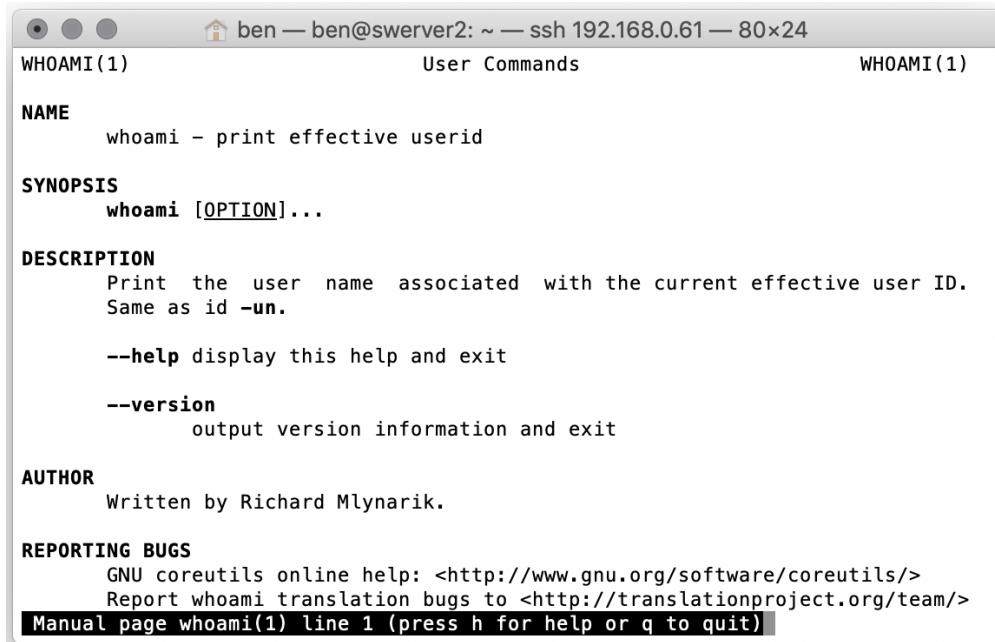
Built-in manuals (the man command)

Linux (and other Unix-like operating systems, such as macOS) have a built-in manual system. It can be accessed in Terminal using the man command. It accepts an argument that is the program or function whose manual page you want to read.

For example, to read information about the whoami command, use the following:

```
man whoami
```

The following figure shows this code running in Terminal:



A screenshot of a terminal window titled "ben — ben@swwerver2: ~ — ssh 192.168.0.61 — 80x24". The window displays the man page for the "whoami" command. The page includes sections for NAME, SYNOPSIS, DESCRIPTION, AUTHOR, and REPORTING BUGS. The DESCRIPTION section notes that "whoami" prints the user name associated with the current effective user ID, and "Same as id -un.". The AUTHOR section credits Richard Mlynarik. The REPORTING BUGS section provides links to online help and a translation bugs report. At the bottom of the page, a message reads "Manual page whoami(1) line 1 (press h for help or q to quit)".

```
WHOAMI(1) User Commands WHOAMI(1)

NAME
    whoami - print effective userid

SYNOPSIS
    whoami [OPTION]...

DESCRIPTION
    Print the user name associated with the current effective user ID.
    Same as id -un.
    --help display this help and exit
    --version
        output version information and exit

AUTHOR
    Written by Richard Mlynarik.

REPORTING BUGS
    GNU coreutils online help: <http://www.gnu.org/software/coreutils/>
    Report whoami translation bugs to <http://translationproject.org/team/>
    Manual page whoami(1) line 1 (press h for help or q to quit)
```

Figure 17.47: A screenshot of man whoami running in Terminal

The information that is displayed is known as a manual page, or “man page.” *Figure 17.47* shows the man page for whoami. You can even run `man man` to read the man page for `man` itself.

SSH

SSH, which stands for **secure shell**, is a protocol to connect to a remote computer and run commands on it from a local system. Using the Terminal interface we’ve worked with in this chapter, SSH allows us to run commands on a remote server, such as a virtual machine or a virtual server in the cloud.

So far, when working with our virtual machine, we have done so using the virtual screen. This is fine to run simple commands, but the virtual screen has some limitations. Copying and pasting, text-selecting, and scrolling all don’t work very well. We also can only have one visible session at a time, which can slow down our workflow. When we begin to work with a virtual server, the virtual screen may be even more difficult to work with, and our provider may offer limited functionality. For these reasons, we will use SSH to communicate with our virtual machine from now on. We already set up an SSH server on the virtual machine when we installed Ubuntu; we just need to get an SSH client set up and connected to it. On macOS and Linux, SSH is built in. On Windows, we will have to install an SSH client. We will use an SSH client called **PuTTY**.

VirtualBox offers different methods of networking, which you can change depending on what works for your situation. When our virtual machine was set up, VirtualBox defaulted to use the **Network Address Translation (NAT)** networking type. This allows connections out from the virtual machine but does not allow new connections in. This is the default because it is most likely to work in all configurations.

Since we need to connect to the virtual machine, we will need to make some changes to the network settings. The easiest method is to switch the virtual machine to use the **Bridged Adapter** setting. This will make the virtual machine behave as if it were just another computer on our network. It will have its own IP address that is separate from our computer. The only issue with this approach is that your network must have its own DHCP server, which will automatically assign an IP address to the computers that are attached. If it does, then this is the preferred method, as it will most closely mimic the setup of a remote virtual server.

The other option is to leave the network in NAT mode and forward the ports we want to use to the virtual machine. For example, SSH runs on port 22 and HTTP on port 80, so both of these ports should be forwarded to the virtual machine. Then, any connections to a host computer’s IP address on those ports will be forwarded to the virtual machine. A few more steps are required to set this up, and it can interfere with those ports if we already have those services running on our host machine.

In the next exercise, you will make these settings changes and then test whether you can connect to your SSH server running on the virtual machine from your host machine. If you’re familiar with how your virtual server’s network is configured already or use different virtualization software or a server from a cloud provider, you can skip this exercise.

Exercise 17.03 – VirtualBox networking configuration

In this exercise, you will make changes to the VirtualBox configuration in order to get its networking set up. This will allow you to connect to the virtual machine using SSH and, later, with a web browser when you run Django on it. Depending on your operating system and which networking settings work, you may be able to skip some steps. At the end of the exercise, you will know how to connect to your virtual machine using SSH:

1. You first need to make sure you have an SSH client installed. On macOS or Linux, you should already have one installed with the system. To confirm it, open Terminal, type `ssh`, and then press *Enter*. You should see output like the following:

```
$ ssh
usage: ssh [-46AaCfGgKkMNnqsTtVvXxYy] [-B bind_interface]
            [-b bind_address] [-c cipher_spec] [-D [bind_
            address:]port]
            [-E log_file] [-e escape_char] [-F configfile] [-I
            pkcs11]
            [-i identity_file] [-J [user@]host[:port]] [-L
            address]
            [-l login_name] [-m mac_spec] [-O ctl_cmd] [-o
            option] [-p port]
            [-Q query_option] [-R address] [-S ctl_path] [-W
            host:port]
            [-w local_tun[:remote_tun]] destination [command]
```

This shows all the options available to configure how SSH connects.

Recent versions of Windows 10 and 11 include an SSH client and server. This can be run directly from Command Prompt.

However, many Windows users continue to use the PuTTY terminal emulator, which has been available since 1999. If you are on Windows, you can download and install PuTTY from <https://www.chiark.greenend.org.uk/~sgtatham/putty/latest.html>.

Now, we can make changes to the VirtualBox configuration. Launch VirtualBox if it's not already running. You should see the **Oracle VM VirtualBox** window with the Bookr virtual machine selected in the left panel. The virtual machine should still be shut down from the previous exercise.

2. In the right panel, click the **Network** header to open the network settings.
3. In the network settings window that opens, click the **Attached to** menu, which should currently be set to **NAT**. Select **Bridged Adapter** (*Figure 17.48*):

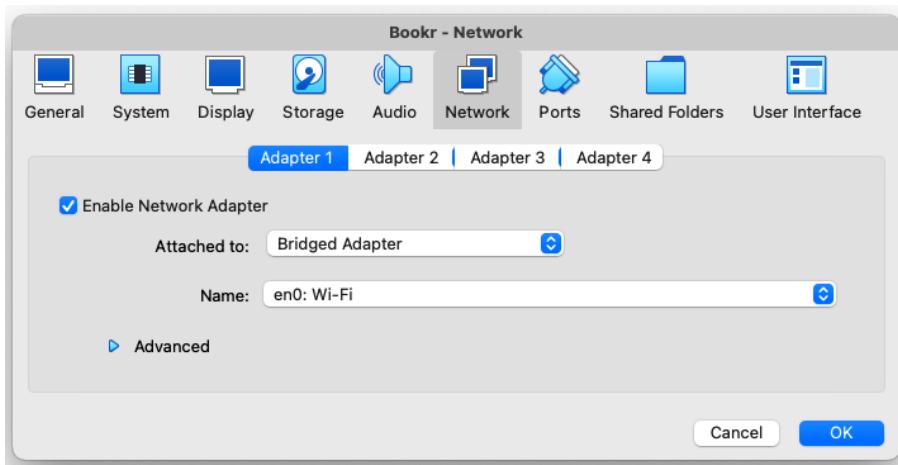


Figure 17.48: Bridged Adapter and the host interface selected

The settings will update to now also show a **Name** field, which allows you to select the physical network device to bridge to.

If your computer has multiple network adapters, select the correct one from this menu. Then, click **OK**.

The virtual machine manager window will be shown again, and the **Network** settings will update to show that it is now in **Bridged Adapter** mode:

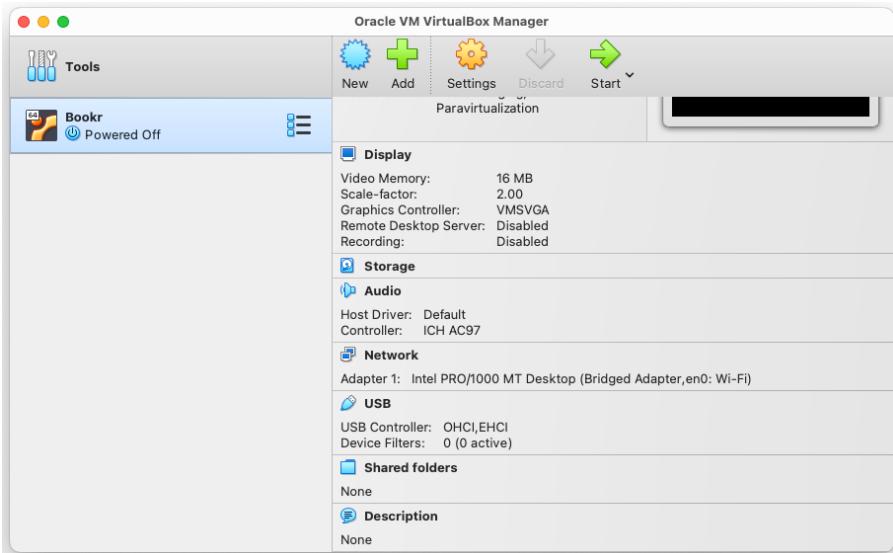
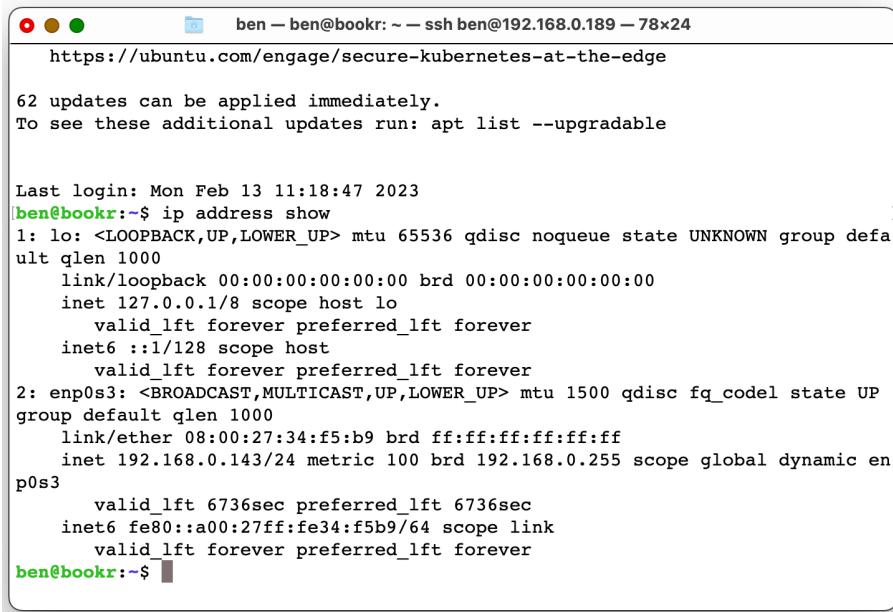


Figure 17.49: The network settings displaying Bridged Adapter

4. Click **Start** at the top of the window to start the virtual machine.
5. After the virtual machine starts up, log in to it using your username and password, as you did in *Exercise 17.02 – Ubuntu Linux installation*. You’ll need to find out the IP address that the virtual machine now has. When you have the Command Prompt, run the `ip` command to see the virtual machine’s IP address.
6. Type `ip address show` and then press *Enter*. You should see output like that shown in *Figure 17.50*:



The screenshot shows a terminal window titled "ben — ben@bookr: ~ — ssh ben@192.168.0.189 — 78x24". The URL "https://ubuntu.com/engage/secure-kubernetes-at-the-edge" is visible above the terminal area. The terminal output is as follows:

```
62 updates can be applied immediately.  
To see these additional updates run: apt list --upgradable  
  
Last login: Mon Feb 13 11:18:47 2023  
[ben@bookr:~]$ ip address show  
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default  
    qlen 1000  
        link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00  
        inet 127.0.0.1/8 scope host lo  
            valid_lft forever preferred_lft forever  
            inet6 ::1/128 scope host  
                valid_lft forever preferred_lft forever  
2: enp0s3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP  
group default qlen 1000  
    link/ether 08:00:27:34:f5:b9 brd ff:ff:ff:ff:ff:ff  
    inet 192.168.0.143/24 metric 100 brd 192.168.0.255 scope global dynamic enp0s3  
        valid_lft 6736sec preferred_lft 6736sec  
        inet6 fe80::a00:27ff:fe34:f5b9/64 scope link  
            valid_lft forever preferred_lft forever  
ben@bookr:~$
```

Figure 17.50: The `ip address show` output

The command lists the network interfaces defined on the virtual machine. The first one, `lo`, refers to the loopback interface. The IP address is listed to the right of `inet` on the third line of the output on the broadcast interface. It should be something such as `192.168.x.y`, `10.x.y.z`, or `172.x.y.z`. If you do not see any of these listed, then it is possible that your computer’s connection is not configured with DHCP, and you will need to use NAT with port forwarding. You can skip forward to step 9, where you’ll configure NAT.

7. If your virtual machine does have an IP address, you can now try connecting to it using SSH. On macOS or Linux, switch to Terminal and run the following:

```
ssh <username>@<ip address>
```

Here, `<username>` is the short username you used when setting up the virtual machine, and `<ip address>` is the IP address of your virtual machine that you got in the previous step. For the virtual machine setup in this example chapter, it would be the following:

```
ssh ben@192.168.0.123
```

Note

Note that the username (highlighted) does not need to be provided if the username on your host machine matches that of your virtual machine. If so, you can just use `ssh <ip address>` – for example, `ssh 192.168.0.123`.

After you press *Enter*, ssh will try to connect to the virtual machine. You will be asked whether you trust the private key from the server, which is displayed as follows:



The screenshot shows a terminal window titled "ben — ben@BensMBP: ~ — ssh ben@192.168.0.123 — 80x24". The window contains the following text:

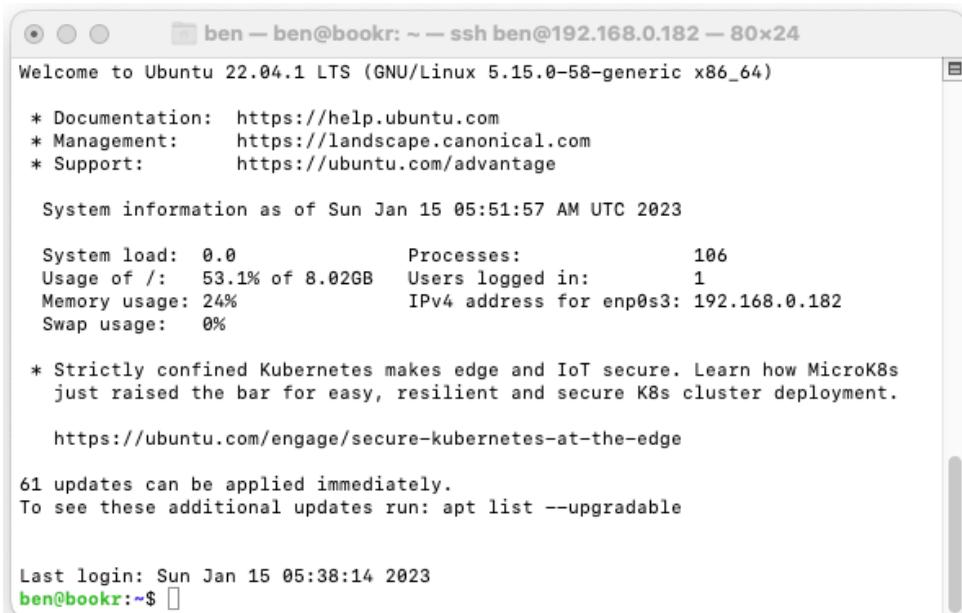
```
ben@BensMBP:~$ ssh ben@192.168.0.123
The authenticity of host '192.168.0.123 (192.168.0.123)' can't be established.
ECDSA key fingerprint is SHA256:tsd6a66lzh2C0hi16WxGLMRrgomPR+YbNqPk6MDn2U.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added '192.168.0.123' (ECDSA) to the list of known hosts.
ben@192.168.0.123's password: ?
```

Figure 17.51: The server private key trust prompt

This is only done the first time you connect to your virtual machine on this IP address (if the IP address changes, you will be prompted again).

8. Type yes, and then press *Enter*.

You will be prompted for your password, which you created in the previous exercise (*step 14* in *Exercise 17.02 – Ubuntu Linux installation*). Type it in (you won't see any output while you type), and then press *Enter* (*Figure 17.52*):



```
ben — ben@bookr: ~ — ssh ben@192.168.0.182 — 80x24
Welcome to Ubuntu 22.04.1 LTS (GNU/Linux 5.15.0-58-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:     https://landscape.canonical.com
 * Support:        https://ubuntu.com/advantage

System information as of Sun Jan 15 05:51:57 AM UTC 2023

System load:  0.0          Processes:           106
Usage of /:   53.1% of 8.02GB  Users logged in:      1
Memory usage: 24%          IPv4 address for enp0s3: 192.168.0.182
Swap usage:   0%

* Strictly confined Kubernetes makes edge and IoT secure. Learn how MicroK8s
just raised the bar for easy, resilient and secure K8s cluster deployment.

https://ubuntu.com/engage/secure-kubernetes-at-the-edge

61 updates can be applied immediately.
To see these additional updates run: apt list --upgradable

Last login: Sun Jan 15 05:38:14 2023
ben@bookr:~$
```

Figure 17.52: After a successful SSH connection

After logging in successfully, you will see the same output as when you logged in using the virtual screen. Then, you'll get a Command Prompt.

9. If you're on Windows, you may prefer to use PuTTY. When you open it, you'll see the **PuTTY Configuration** window. Enter the IP address you found in *step 4* into the **Host Name (or IP address)** field. Then, click **Open** at the bottom of the window:

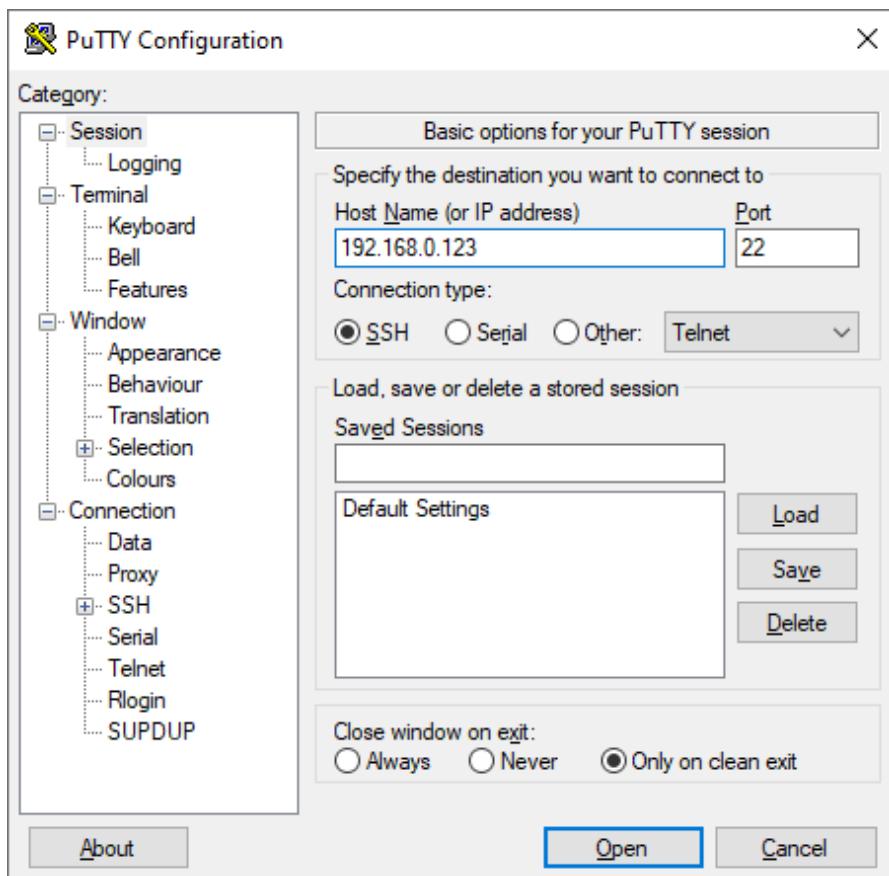


Figure 17.53: The PuTTY Configuration window

10. Since this is the first time connecting to the virtual machine, you will see the **PuTTY Security Alert** window, which displays the server's key:

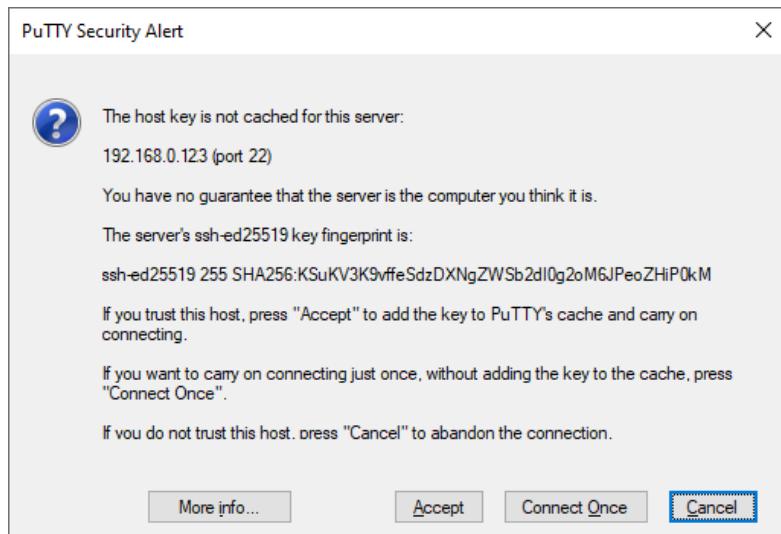


Figure 17.54: The PuTTY Security Alert window

11. Click **Accept**.

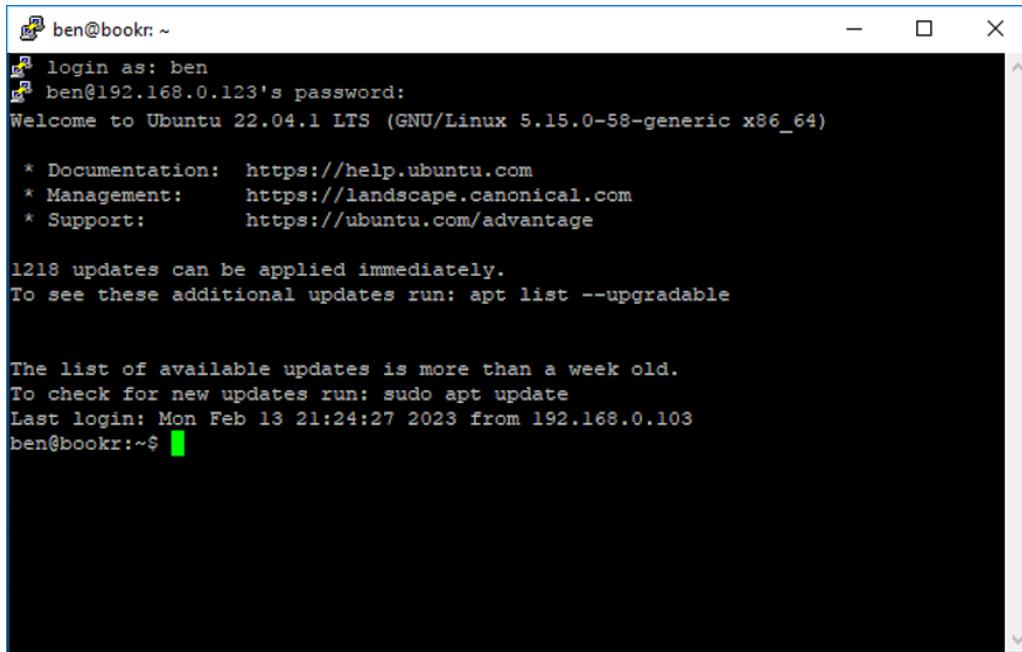
You'll then see a Terminal window, which will first prompt you who to log in as. Enter the username you created in the previous exercise, and then press *Enter*.

It will then prompt you for your password. Type your password (you won't see anything on screen as you type), and then press *Enter* again:



Figure 17.55: The PuTTY login/password prompt

After logging in successfully, you will see the same output as when you logged in using the virtual screen. Then, you'll get a Command Prompt, as shown in the following figure:



```
ben@bookr: ~
└─ ben
└─ ben@192.168.0.123's password:
Welcome to Ubuntu 22.04.1 LTS (GNU/Linux 5.15.0-58-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:     https://landscape.canonical.com
 * Support:        https://ubuntu.com/advantage

1218 updates can be applied immediately.
To see these additional updates run: apt list --upgradable

The list of available updates is more than a week old.
To check for new updates run: sudo apt update
Last login: Mon Feb 13 21:24:27 2023 from 192.168.0.103
ben@bookr:~$
```

Figure 17.56: Command Prompt after logging in using PuTTY

Regardless of your operating system, if you were able to log in and see the Command Prompt, this means everything is working fine. Make sure you record the IP address of the virtual machine because you'll need it every time you have to log in to the server, but for now, we know that it works. You can skip ahead to the final step, *step 17*.

12. If, for some reason, the virtual machine did not get an IP address or the SSH connection failed, you will need to try to connect using NAT. Switch back to the VirtualBox virtual screen, where you should still be at the Command Prompt where you ran the `ifconfig` command. Shut down the virtual machine by running the `shutdown now` command.
13. In the **Oracle VM VirtualBox Manager** window, click the **Network** heading text (as you did in *step 2*).

14. In the **Network** settings window, change the **Attached to** setting back to **NAT** (not **NAT Network**):

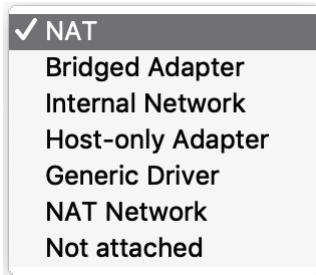


Figure 17.57: Attached to changed back to NAT

15. You now need to forward ports so that connections to your computer will go through to the virtual machine. Click the **Advanced** arrow to expand the advanced settings section:

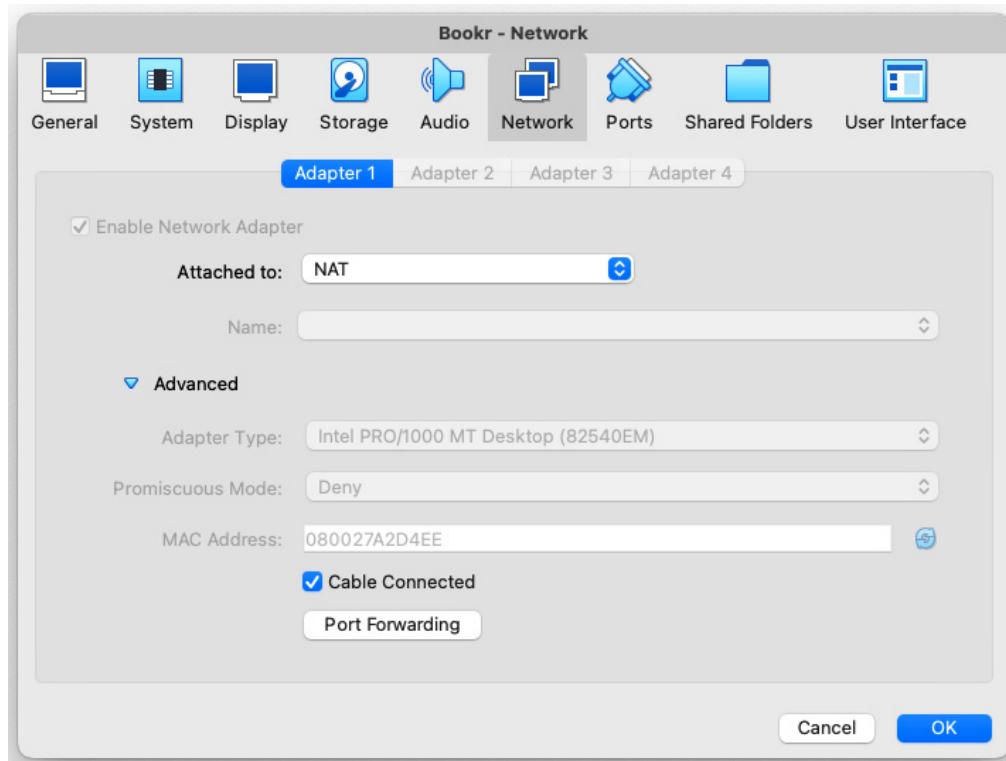


Figure 17.58: The advanced network settings opened

16. Click the **Port Forwarding** button to open the port forwarding settings. Then, click the **Add** button in the top-right corner. You will see a blank port forwarding configuration row displayed.

Refer to *Figure 17.59* of step 17 to see what this window looks like. Enter the following values:

- **Name:** SSH.
- **Protocol:** Select TCP.010
- **Host Port:** 22.
- **Guest Port:** 22.

The other fields (**Host IP** and **Guest IP**) should be left blank.

This will enable SSH to be forwarded.

17. Repeat *step 16* to add another port forwarding line. This time, enter HTTP for **Name**, and 80 for both **Host Port** and **Guest Port**. When you've completed *step 16* and this step, the port forwarding settings should look like *Figure 17.59*:

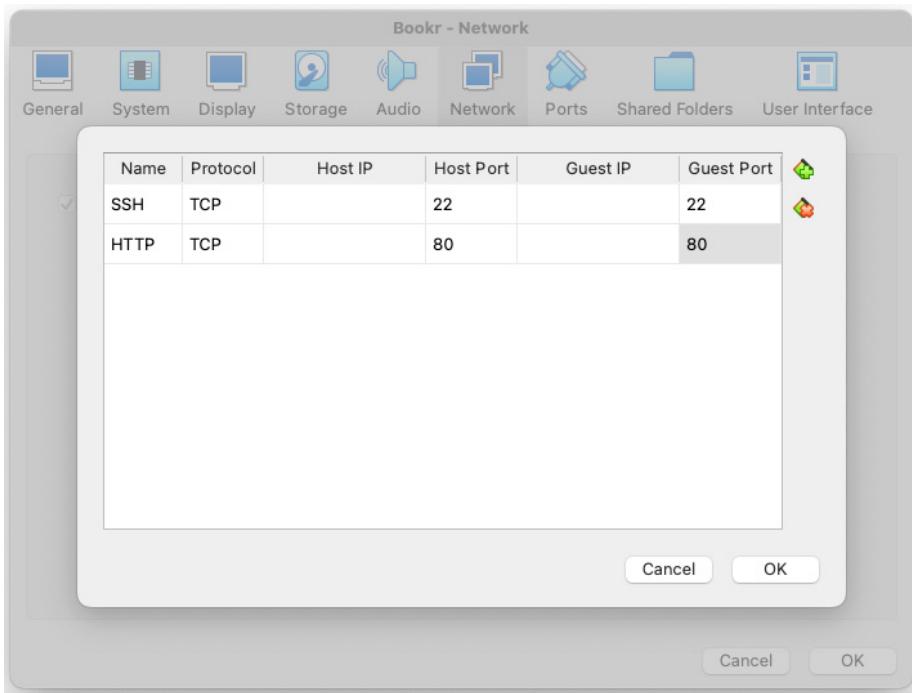


Figure 17.59: The port forwarding settings

18. Click **OK** to close and save the port forwarding settings, and then click **OK** in the **Network** settings window to save and close it.

19. Start the virtual machine by clicking **Start** in the **Oracle VM VirtualBox Manager** window.
20. Once it has started up, you don't need to log in to the virtual screen; instead, switch back to Terminal (macOS/Linux) or PuTTY (Windows) and try to connect (repeat *step 7* or *step 9*) using the `localhost` address. This won't connect to your host computer; since the port is forwarded, it will go to the virtual machine instead.

You should now be logged in and see the Command Prompt, as shown in *Figure 17.52* (macOS/Linux) and *Figure 17.56* (Windows).

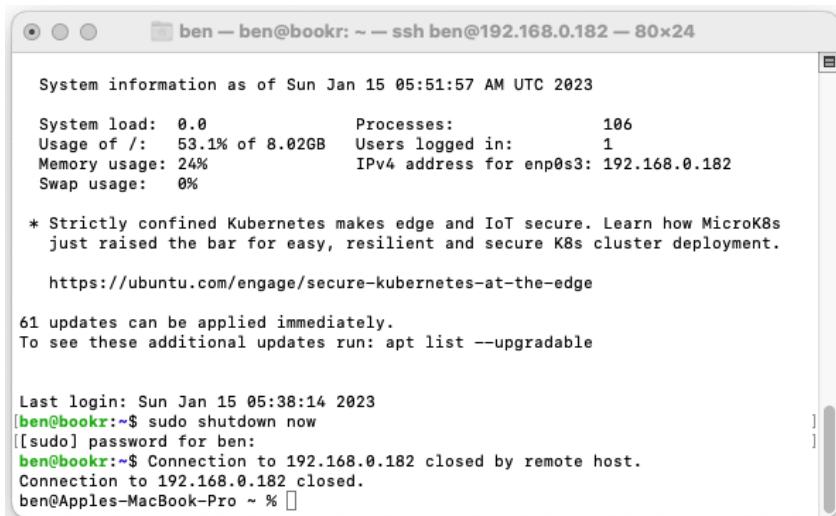
21. You can now try out the `sudo` command to shut down the virtual machine.

Note

When we used the virtual screen inside VirtualBox, we could run the `shutdown` command without using `sudo`. This is because it was like we were sitting down in front of a real computer and using a keyboard and monitor connected to it. Anyone in this position can run `shutdown` because even if they did not have software access, they would still have hardware access to just turn the computer off with a power switch.

When we connect over SSH, we are not physically at the machine, so the `shutdown` command is limited to `root`.

22. Shut down your virtual machine now by typing `sudo shutdown now`, and then press *Enter*. You will be prompted to enter your password. Type it in (nothing will be shown on screen), and then press *Enter*:



The screenshot shows a terminal window titled "ben — ben@bookr: ~ — ssh ben@192.168.0.182 — 80x24". The window displays system information and a message about Kubernetes security. It then shows the command "sudo shutdown now" being entered, followed by a password prompt, and finally a message indicating the connection was closed by the remote host.

```
ben — ben@bookr: ~ — ssh ben@192.168.0.182 — 80x24

System information as of Sun Jan 15 05:51:57 AM UTC 2023

System load: 0.0          Processes:          106
Usage of /: 53.1% of 8.02GB  Users logged in:  1
Memory usage: 24%
Swap usage:  0%

* Strictly confined Kubernetes makes edge and IoT secure. Learn how MicroK8s
just raised the bar for easy, resilient and secure K8s cluster deployment.

https://ubuntu.com/engage/secure-kubernetes-at-the-edge

61 updates can be applied immediately.
To see these additional updates run: apt list --upgradable

Last login: Sun Jan 15 05:38:14 2023
[ben@bookr:~$ sudo shutdown now
[[sudo] password for ben:
ben@bookr:~$ Connection to 192.168.0.182 closed by remote host.
Connection to 192.168.0.182 closed.
ben@Apples-MacBook-Pro ~ % ]]
```

Figure 17.60: The shutdown command executed over SSH

The virtual machine will shut down and your SSH session will be disconnected.

In this exercise, you configured networking for a virtual machine and used SSH to connect to it. Your virtual machine may use either bridged or NAT networking, depending on which setting worked.

If you set up a virtual server with a cloud provider, it would be in this current state, so the steps and exercises going forward apply to both your local virtual machine and the remote virtual private server hosted in the cloud.

In the next section, you will learn how to install packages and keep your system up to date using `apt`. You'll understand the various system packages that you need and what they do. Then, you'll log back in to your virtual machine and install them.

Packages and updates

Different distributions of Linux have different systems to install and update packages. Ubuntu (and other distributions that are also based on Debian) uses a system called **Advanced Package Tool (APT)**. This is a system to install and update packages, as well as automatically install dependencies and default configuration files. It is also able to run pre- and post-installation scripts to perform an extra setup.

There are a few different ways to interact with the APT system; we will use the `apt` command-line tool.

Since `apt` needs to make changes to the system, it must be run using `sudo`. Here's a quick rundown of a few common commands:

- `install`: This is used to install a particular package:

```
sudo apt install <packagename>
```

This will also first install the packages on which the installed package(s) depend. You can also install multiple packages at once, by listing them all – for example, `sudo apt install package1 package2 ...`.

- `remove`: This will uninstall a package:

```
sudo apt remove <packagename>
```

It can take multiple package names, as the `install` command does. This will not remove the package's dependencies (see `autoremove`).

- `purge`: When `apt` removes a package (using the `remove` command), it won't delete any configuration files that the package installed since you might have made changes to them that you don't want to lose. Alternatively, you might need to remove an old version of a package and reinstall a new one that uses the same config files. Whatever the reason, if you do also want to clean up everything, `apt` provides the `purge` command. `purge` is used in the same way as `remove` but also deletes configuration files:

```
sudo apt purge <packagename>
```

- **update:** This refreshes the list of available packages, including new packages and upgrades. It does not take any arguments:

```
sudo apt update
```

The package list must be updated before trying to upgrade packages; otherwise, the local system will not know that upgrades exist.

- **upgrade:** This command upgrades all the installed packages to their latest version:

```
sudo apt upgrade
```

This does not take any arguments – that is, you can't use it to upgrade just one installed package. Instead, just `apt install` a package again, and only that package will be upgraded.

- **full-upgrade:** This is similar to `upgrade`; however, when using `upgrade`, new dependencies are not installed. For example, you might install `package-a`, which is version 1.0 and has no dependencies. Later, a new version, 2.0, is released, which now has a dependency, `package-b`. Running `upgrade` will not upgrade `package-a` because it must install the `package-b` dependency. Running `full-upgrade` will first install `package-b`, and then version 2.0 of `package-a`.

Like `upgrade`, it takes no arguments:

```
sudo apt full-upgrade
```

- **autoremove:** When `apt` removes a package (using the `remove` command), it leaves that package's dependencies installed. The `autoremove` command will remove all the packages that have been installed as dependencies but whose dependent packages have been removed.

This command takes no arguments:

```
sudo apt autoremove
```

In this book, we will only use the `update` and `install` commands, although knowing the others that have been covered will give you a good basis to continue to administer the Ubuntu system.

Ubuntu comes with some of the software that we need already installed – for example, Python 3. To run a Django app, we'll need some other supporting packages:

- **nginx:** This package installs the NGINX web server to serve static files and proxy requests to Gunicorn.
- **postgresql:** This is the SQL server that replaces the single-file SQLite database we used during development.
- **python3-virtualenv:** This adds the Python `virtualenv` package to the system. Once this is installed, we'll be able to create Python virtual environments. Further Python packages (such as Django) will be installed into the virtual environment using `pip`.

- `libpq-dev`, `python3-dev`, and `build-essential`: These provide header files for PostgreSQL, Python 3, and the system. `build-essential` also includes compilers and build tools. We'll need these to build the Python PostgreSQL client library.

In the next exercise, we will upgrade the packages that are already installed on the virtual machine. Then, we'll install the packages we just mentioned.

Exercise 17.04 – package update and installation

In this exercise, you will SSH into your virtual machine, and then update the list of available packages/versions using the `apt update` command. The installed packages will then be upgraded using `apt upgrade`. Finally, we'll use `apt install` to install `nginx`, `postgresql`, `python3-virtualenv`, `libpq-dev`, `python3-dev`, and `build-essential`:

1. Start VirtualBox and then start up your virtual machine. Let it finish booting up – you'll know it's finished because you'll see the login prompt.
2. SSH into your virtual machine, using the IP address you found in *Exercise 17.03 – VirtualBox networking configuration* (or `localhost`).
3. Once you have connected, use the `apt update` command to update the list of available packages. The command must be run using `sudo`, like this:

```
sudo apt update
```

You will be prompted to enter your password to authenticate the `sudo` process:

```
ben@bookr:~$ sudo apt update  
[sudo] password for ben:
```

4. Type in your password and then press *Enter*.

You will see the list of repositories from which `apt` fetches updates, and you should see that several packages are available for upgrade:

```
ben@bookr:~$ sudo apt update  
[sudo] password for ben:  
Hit:1 http://nz.archive.ubuntu.com/ubuntu jammy InRelease  
Get:2 http://nz.archive.ubuntu.com/ubuntu jammy-updates  
InRelease  
Get:3 http://nz.archive.ubuntu.com/ubuntu jammy-backports  
InRelease  
Get:4 http://nz.archive.ubuntu.com/ubuntu jammy-security  
InRelease  
...  
Fetched 5,037 kB in 3s (1,560 kB/s)  
Reading package lists... Done  
Building dependency tree... Done  
Reading state information... Done
```

```
77 packages can be upgraded. Run 'apt list --upgradable' to see them.  
ben@bookr:~$
```

5. Upgrade all the installed packages to their latest versions using the `apt upgrade` command. Again, this must be run using `sudo`:

```
sudo apt upgrade
```

Since you just authenticated to `sudo` in the previous step, you shouldn't have to enter your password again (unless it's been more than 5 minutes since you executed *step 3*).

After entering the command, `apt` will check which packages are upgradeable and show them to you. It will then ask the following:

```
Do you want to continue? [Y/n]
```

You'll need to then press *Enter* again to confirm that you want to upgrade. `apt` will then download and install all the available upgrades (note that your list of upgrades might not exactly match what is shown here, and for brevity, a lot of output has been removed):

```
ben@bookr:~$ sudo apt upgrade  
Reading package lists... Done  
Building dependency tree  
Reading state information... Done  
Calculating upgrade... Done  
The following packages will be upgraded:  
  bsdutils distro-info-data ...  
28 upgraded, 0 newly installed, 0 to remove and 0 not upgraded.  
Need to get 7344 kB of archives.  
After this operation, 238 kB of additional disk space will be used.  
Do you want to continue? [Y/n]  
Get:1 http://nz.archive.ubuntu.com/ubuntu bionic-updates/main  
amd64 bsdutils amd64 1:2.31.1-0.4ubuntu3.6 [60.3 kB]  
Get:2 http://nz.archive.ubuntu.com/ubuntu bionic-updates/main  
amd64 libuuid1 amd64 2.31.1-0.4ubuntu3.6 [20.1 kB]  
...  
Fetched 7344 kB in 4s (1720 kB/s)  
Preconfiguring packages ...  
(Reading database ... 102499 files and directories currently  
installed.)  
Preparing to unpack .../bsdutils_1%3a2.31.1-0.4ubuntu3.6_amd64.  
deb ...  
Unpacking bsdutils (1:2.31.1-0.4ubuntu3.6) over (1:2.31.1-  
0.4ubuntu3.5) ...  
Setting up bsdutils (1:2.31.1-0.4ubuntu3.6) ...  
...  
Processing triggers for initramfs-tools (0.130ubuntu3.9) ...
```

```
update-initramfs: Generating /boot/initrd.img-4.15.0-99-generic  
ben@bookr:~$
```

- Now that all the existing packages are up to date, install the packages you need - `nginx`, `postgresql`, `python3-virtualenv`, `libpq-dev`, `build-essential`, and `python3-dev`. This is done using `apt install`, and all the package names can be provided to the command at once:

```
sudo apt install nginx postgresql python3-virtualenv libpq-dev  
build-essential python3-dev
```

`apt` will calculate the package dependencies and show you a list of all the packages that will be installed. Some of the output has been truncated for brevity:

```
ben@bookr:~$ sudo apt install nginx postgresql python3-  
virtualenv libpq-dev build-essential python3-dev  
[sudo] password for ben:  
Reading package lists... Done  
Building dependency tree... Done  
Reading state information... Done  
The following packages were automatically installed and are no  
longer required:  
  libflashrom1 libftdi1-2  
Use 'sudo apt autoremove' to remove them.  
The following additional packages will be installed:  
  bzip2 cpp cpp-11 dpkg-dev fakeroot fontconfig-config ...  
  zlib1g-dev  
Suggested packages:  
  bzip2-doc cpp-doc gcc-11-locales debian-keyring g++-multilib ...  
  python2-pip-whl python2-setuptools-whl isag  
The following NEW packages will be installed:  
  build-essential bzip2 cpp cpp-11 dpkg-dev fakeroot ...  
  python3.10-dev rpcsvc-proto ssl-cert sysstat zlib1g-dev  
0 upgraded, 97 newly installed, 0 to remove and 4 not upgraded.  
Need to get 120 MB of archives.  
After this operation, 420 MB of additional disk space will be  
used.  
Do you want to continue? [Y/n]
```

Once again, you will be asked whether you want to continue, so press `Enter` to begin the installation.

Depending on the speed of your internet connection, the installation might take a few minutes or more. You will see a lot of output. First will be messages starting with `Get`, as follows:

```
Get:1 http://nz.archive.ubuntu.com/ubuntu jammy-updates/main  
amd64 libc-dev-bin amd64 2.35-0ubuntu3.1 [20.4 kB]  
Get:2 http://nz.archive.ubuntu.com/ubuntu jammy-updates/main  
amd64 linux-libc-dev amd64 5.15.0-58.64 [1,347 kB]
```

This indicates the packages that are being fetched. Next, the packages will be unpacked, with lines like this:

```
Selecting previously unselected package libc-dev-bin.  
(Reading database ... 73777 files and directories currently  
installed.)  
Preparing to unpack .../00-libc-dev-bin_2.35-0ubuntu3.1_amd64.  
deb ...  
Unpacking libc-dev-bin (2.35-0ubuntu3.1) ...  
Selecting previously unselected package linux-libc-dev:amd64.  
Preparing to unpack .../01-linux-libc-dev_5.15.0-58.64_amd64.deb  
...  
Unpacking linux-libc-dev:amd64 (5.15.0-58.64) ...
```

Next, the packages are set up:

```
Setting up postgresql-client-common (238) ...  
Setting up javascript-common (11+nmu1) ...  
Setting up gcc-11-base:amd64 (11.3.0-1ubuntu1~22.04) ...
```

For some packages (postgresql in particular), there will be a lot of output about how they are configured. When apt installs nginx and postgresql, it also sets them up to start when the machine boots, and then starts them after the installation is complete.

The last few lines of output involve the processing of triggers, which are the final scripts to set up the packages:

```
Processing triggers for ufw (0.36.1-4build1) ...  
Processing triggers for man-db (2.10.2-1) ...  
Processing triggers for libc-bin (2.35-0ubuntu3.1) ...
```

When the installation is finished, you will see your Command Prompt again.

7. The final step in this exercise is to check that NGINX was correctly installed. Open a web browser on your host machine and go to the address of your virtual machine – for example, <http://192.168.0.123/> or <http://localhost/> (if using NAT). You should see the default NGINX welcome page:

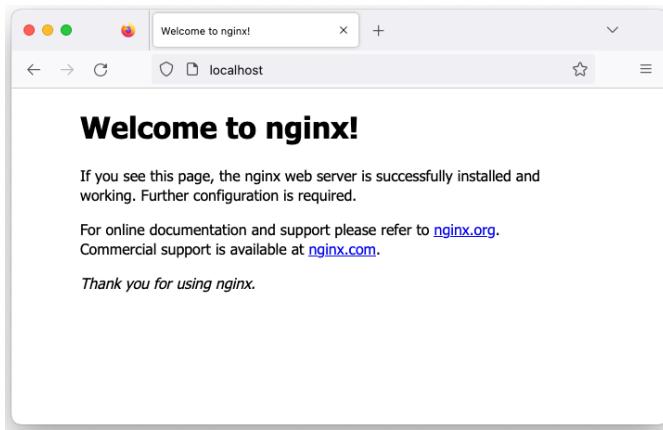


Figure 17.61: The default NGINX screen

It is not so easy to test that PostgreSQL was set up correctly, but unless you saw any errors in the installation output, you should assume it was successful. You'll use PostgreSQL in the first exercise in the next chapter.

8. We will return to the virtual machine in the next chapter. You can shut it down now by running the `sudo shutdown` command in the SSH terminal. Alternatively, if you want to continue to the next chapter straightaway, you can leave the session running and return to it.

In this exercise, you first updated the list of available packages and then installed all the available upgrades. Then, you installed the NGINX web server, the PostgreSQL database server, Python 3's `virtualenv`, and other supporting build packages. This was all done using the `apt` program. You then tested that the NGINX web server was running and that you could connect to it using a web browser.

The following activity will give you an opportunity to use your Linux command line skills to install a useful system monitoring utility.

Activity 17.01 – installing `glances` – a system-monitoring utility

In the activity for this chapter, you will reinforce what you have learned by connecting to your virtual machine and then installing and using the `glances` system monitoring application, all on your own.

`glances` is a Terminal-based utility that shows you information about your virtual machine at a glance (hence the name). It can be installed with `apt` and has several command-line options. To get you more familiar with using SSH, `apt`, and Linux in general, in this activity, you will update your `apt` package list, install `glances`, and then check its command-line options using `man`.

These steps will help you complete this activity:

1. Start up your virtual machine and connect to it using SSH.
2. Make sure its list of `apt` packages is up to date.
3. Install `glances` using `apt`.
4. After it's installed, run `glances` and have a quick look at its interface. Then, press `Ctrl + C` or `Esc` to exit.
5. There are two things you might have noticed when running `glances`.

First, if your Terminal has a light background, some of the text can be difficult to read. `glances` doesn't have a white theme to fix this.

Second, there is a lot of information displayed, and you can disable some of the information that is less interesting to you. In this case, as an example, you want to disable the memory swap module.

Use the `man` command to read about the command-line options for `glances`. You will be able to find out how to disable the memory swap module and, optionally, set it to the white theme (only if your Terminal has a light background).

6. Rerun `glances` with the command-line flags to set the options you learned about in step 5.

Once `glances` is up and running with the correct options, it should look similar to *Figure 17.62* (although it will look different if your Terminal has a dark background):

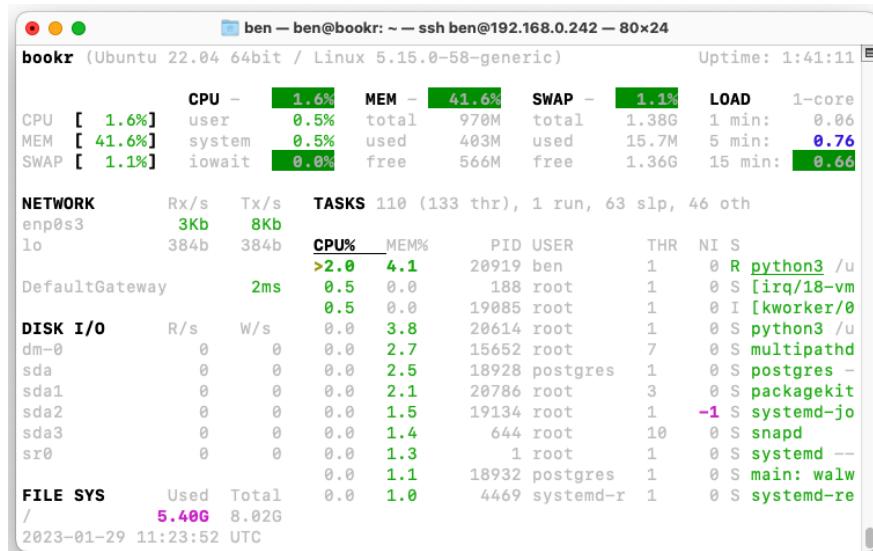


Figure 17.62: `glances` with a white theme and the memory swap module disabled

Note

The solution to this activity can be found at <https://github.com/PacktPublishing/Web-Development-with-Django-Second-Edition/tree/main/ActivitySolutions>.

Summary

In this chapter, we learned about the architecture used in a production web server, specifically the roles of a frontend server (NGINX), an application server (Gunicorn), and databases (PostgreSQL) in Django hosting. We learned how these three work together to receive, process, and serve requests that the server gets.

We then got Ubuntu Linux up and running on a virtual server, either on a personal computer or through a cloud provider. Later, we connected to this server remotely using SSH. Then, we learned about the apt command and how it comes in handy when we need to install or remove packages and keep them up to date. Finally, we put all the knowledge we'd gained so far in the chapter into practice by installing NGINX, Gunicorn, and PostgreSQL, plus the supporting packages, on our virtual server remotely using SSH.

The next chapter follows straight on from this one. We will configure our PostgreSQL database and users, configure permissions on our server, upload our application code and production settings, and finish by getting Bookr up and running on our virtual machine.

